

PAL: A VARIABILITY-AWARE POLICY FOR SCHEDULING ML WORKLOADS IN GPU CLUSTERS

Rutwik Jain, Brandon Tran, Keting Chen, Matthew D. Sinclair, Shivaram Venkataraman

Computer Sciences Department, University of Wisconsin-Madison




Computer Sciences

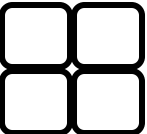
SCHOOL OF COMPUTER, DATA & INFORMATION SCIENCES
UNIVERSITY OF WISCONSIN-MADISON



SC24
Atlanta, GA | hpc creates.

GPU Cluster






 = 1 node (4 GPUs)



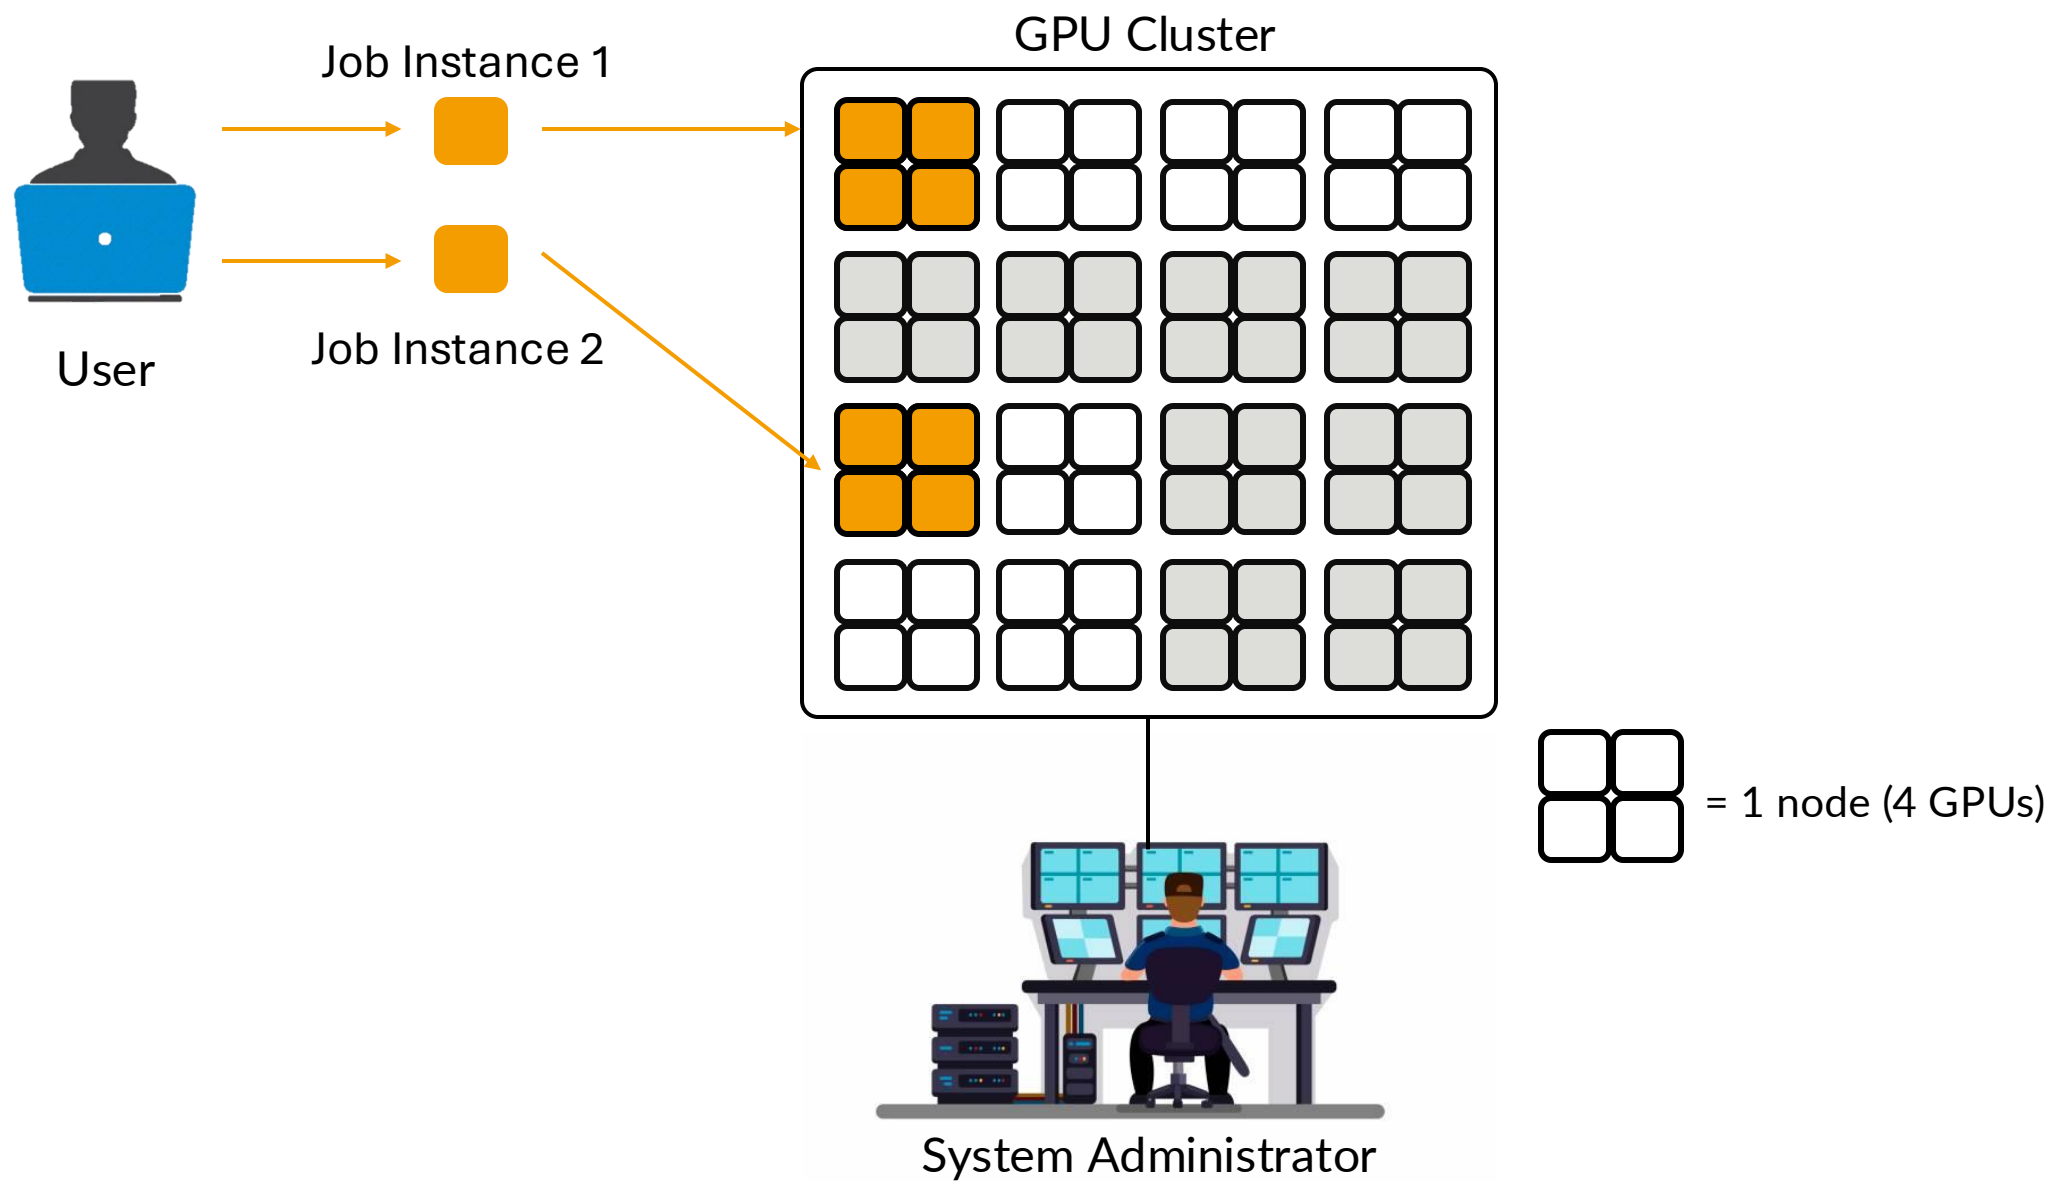
System Administrator

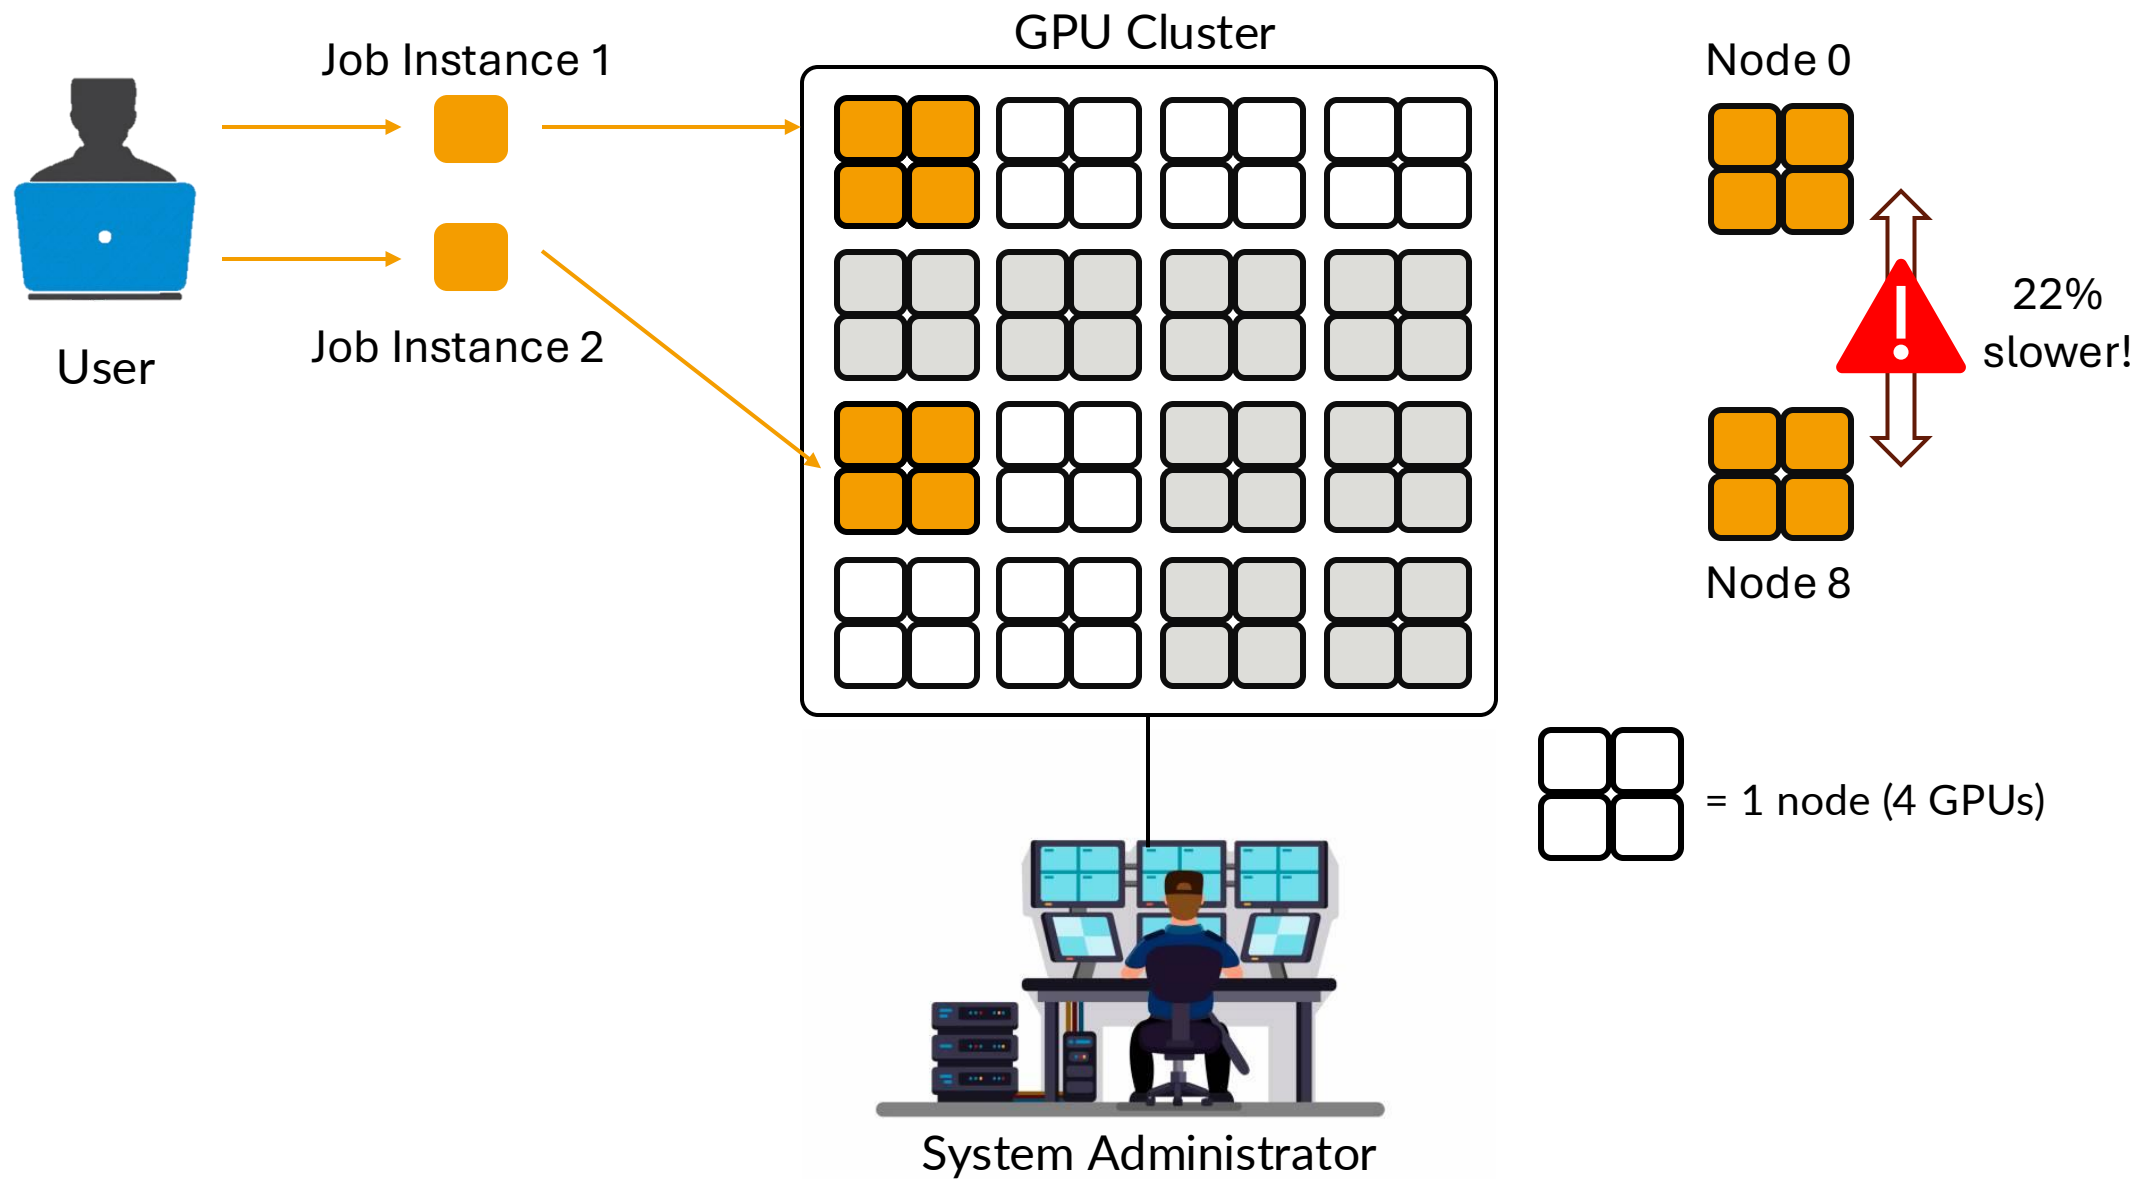




System Administrator





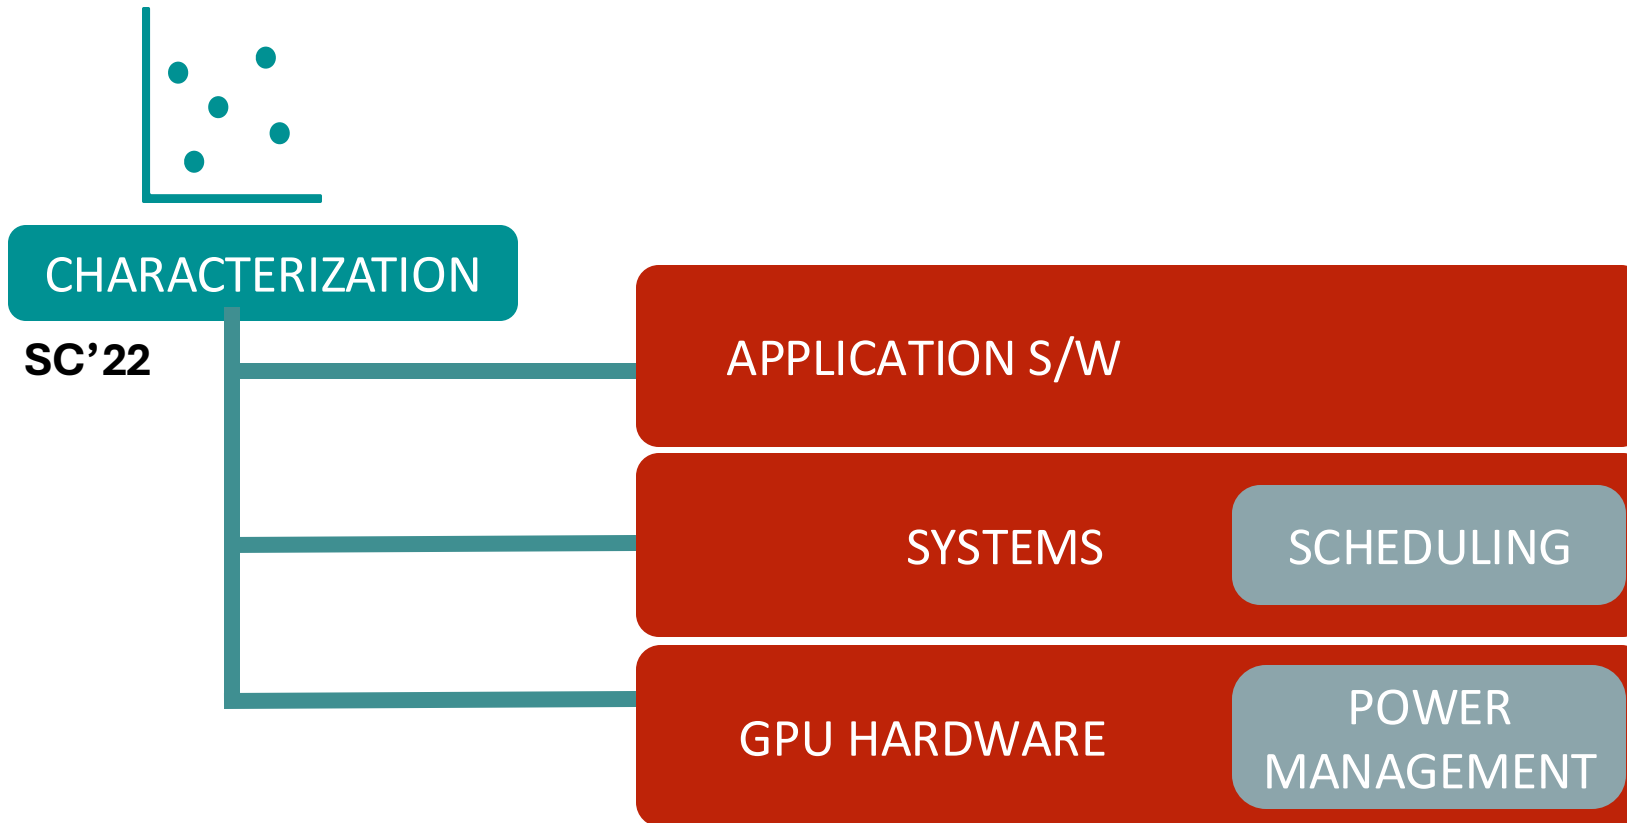


GPU PERFORMANCE VARIABILITY

- Culprit: **Variability!**
- GPUs exhibit variability because of:
 - Local Power Management (PM)
 - Thermal/non-uniform cooling effects
 - Process variation
- **Impact**
 - Users: Hard to get **repeatable, consistent** performance for applications.
 - Sysadmins: **Resource underutilization** for multi-GPU jobs – faster GPUs wait for stragglers



CHARACTERIZING VARIABILITY

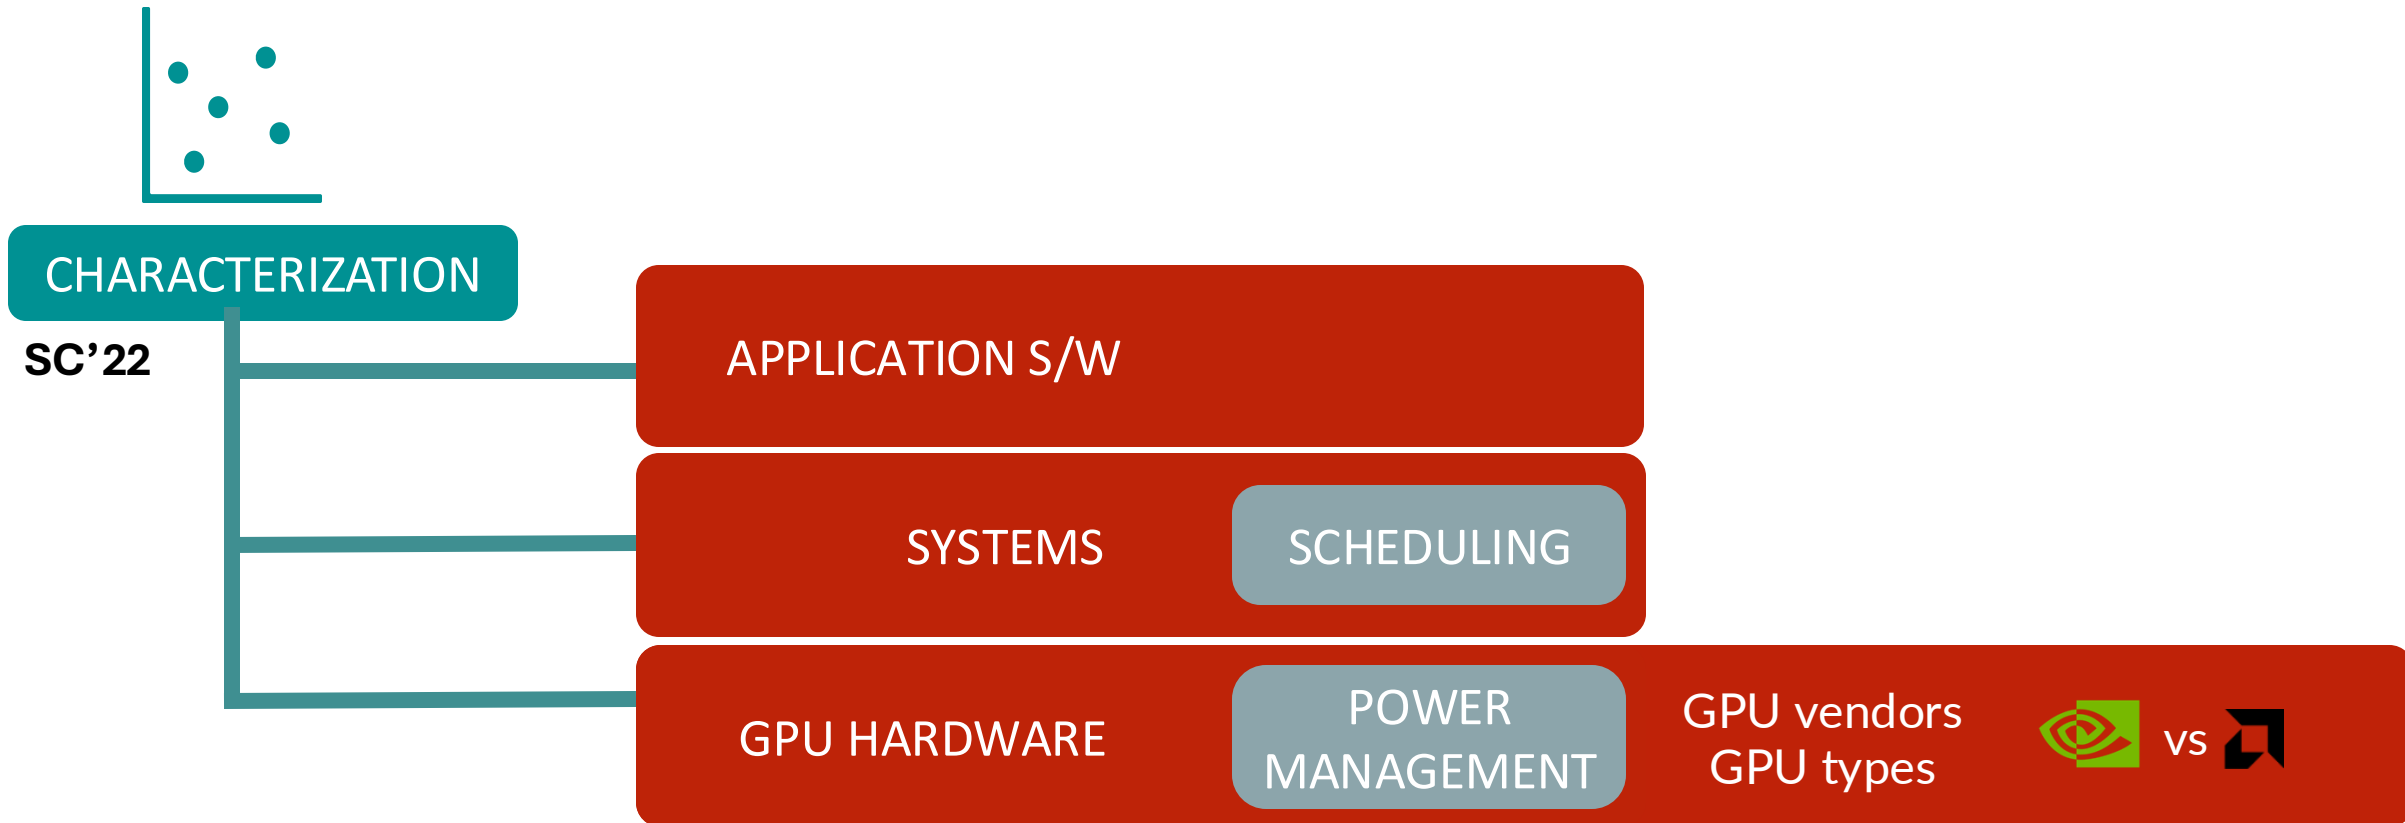


Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



CHARACTERIZING VARIABILITY

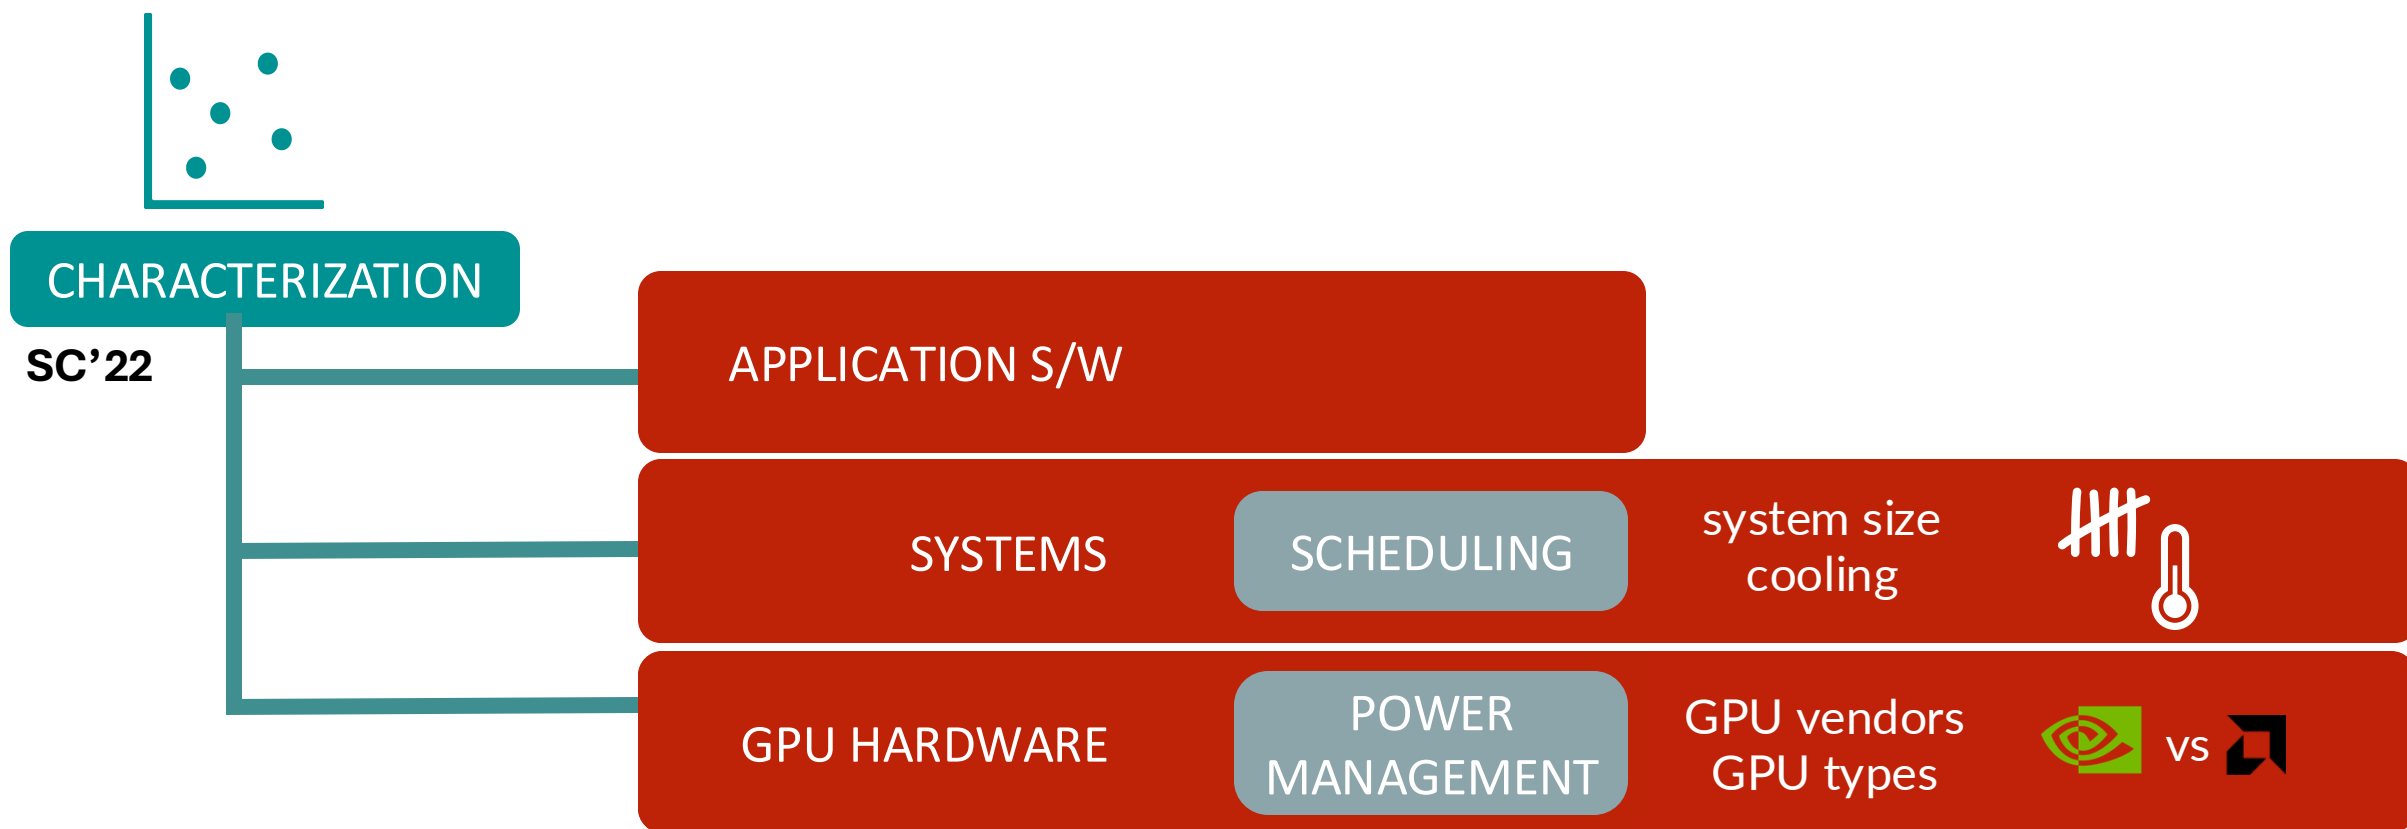


Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



CHARACTERIZING VARIABILITY

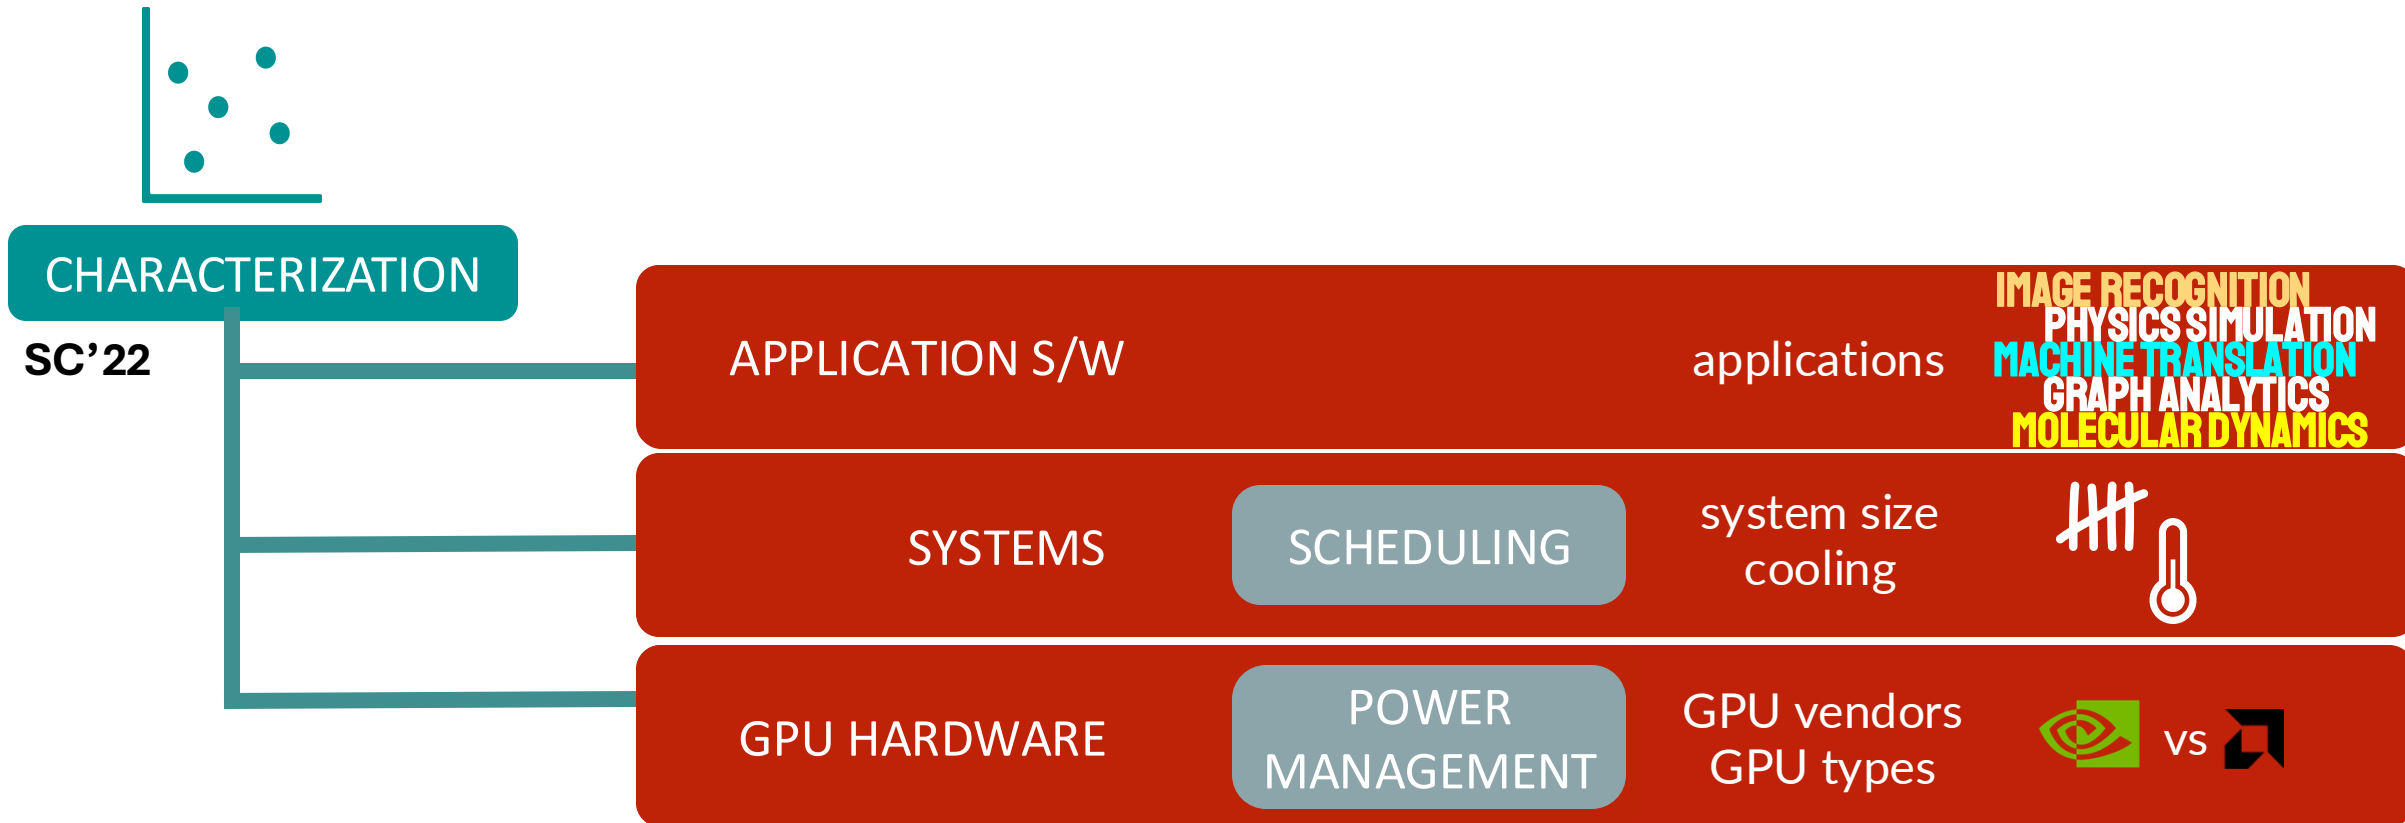


Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



CHARACTERIZING VARIABILITY



Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>

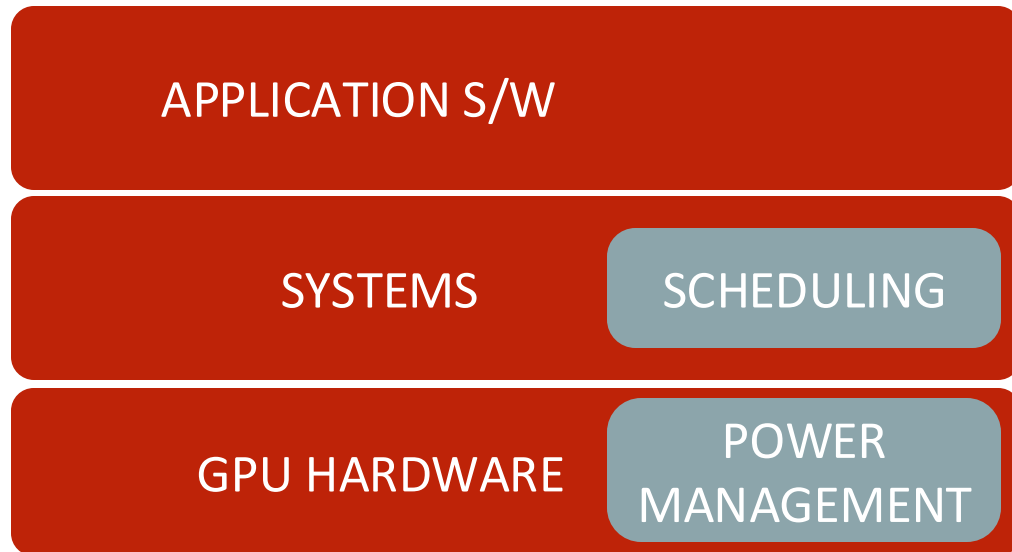


CHARACTERIZATION STUDY: TAKEAWAYS

- 1 Consistent performance variability across clusters, GPU vendors and cooling methods (8% for SGEMM, outliers up to 1.5x slower)
- 2 Variability is not a transient effect, ill-performing GPUs are consistently ill-performing
- 3 **Variability is application-specific** – compute-intensive workloads are more variability sensitive (ResNet-50 shows 22% variability while PageRank has <1%)



KEY INSIGHT: EMBRACE VARIABILITY

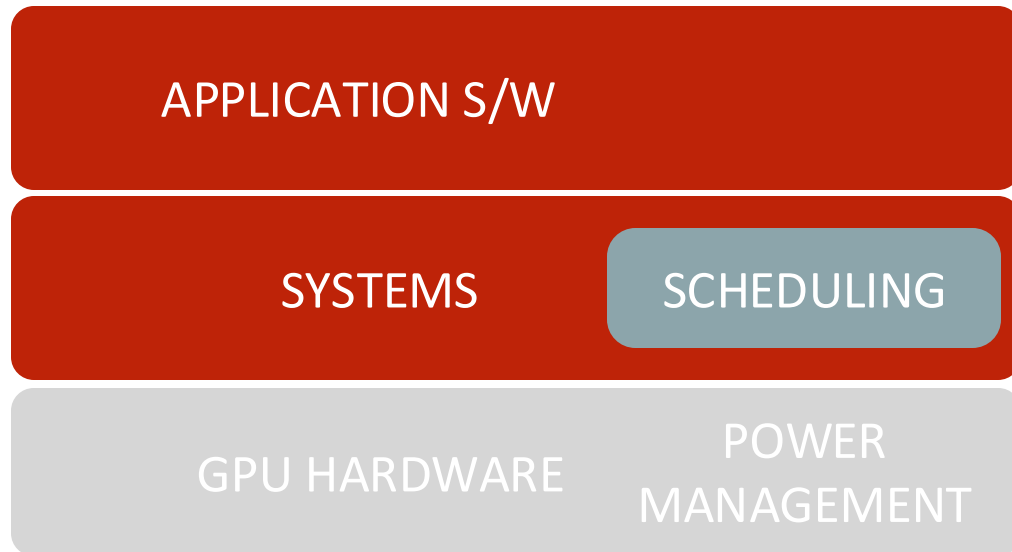


Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



KEY INSIGHT: EMBRACE VARIABILITY

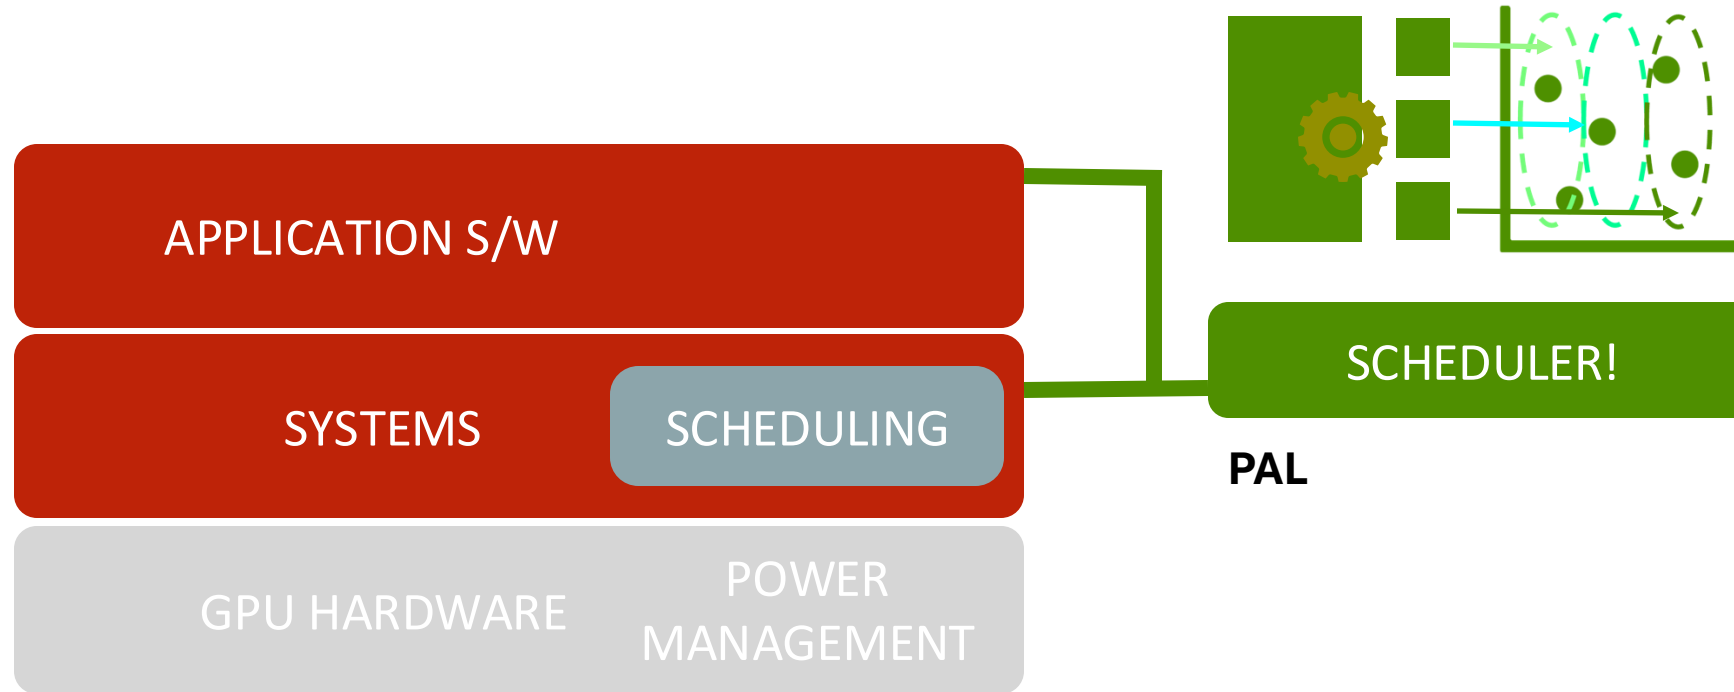


Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



KEY INSIGHT: EMBRACE VARIABILITY



Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems (SC22)

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



OUR APPROACH

- We characterize which applications are more likely to suffer from performance **variability** & take that into account **while placing jobs on the cluster**.
- Novel placement policies: **PM-First** and **PAL**
 - **PM-First** uses **application-specific variability profiles** to improve performance and utilization.
 - **PAL** further improves scheduling by balancing **variability** with **locality**.
- Overall, PAL improves geomean Job Completion Time (JCT) by **42%** and cluster utilization by **28%** over state-of-the-art schedulers like Tiresias




OUTLINE


- Introduction
- **Cluster Scheduling**
- **Design**
- **Methodology**
- **Evaluation**
- **Conclusion**



CLUSTER SCHEDULING



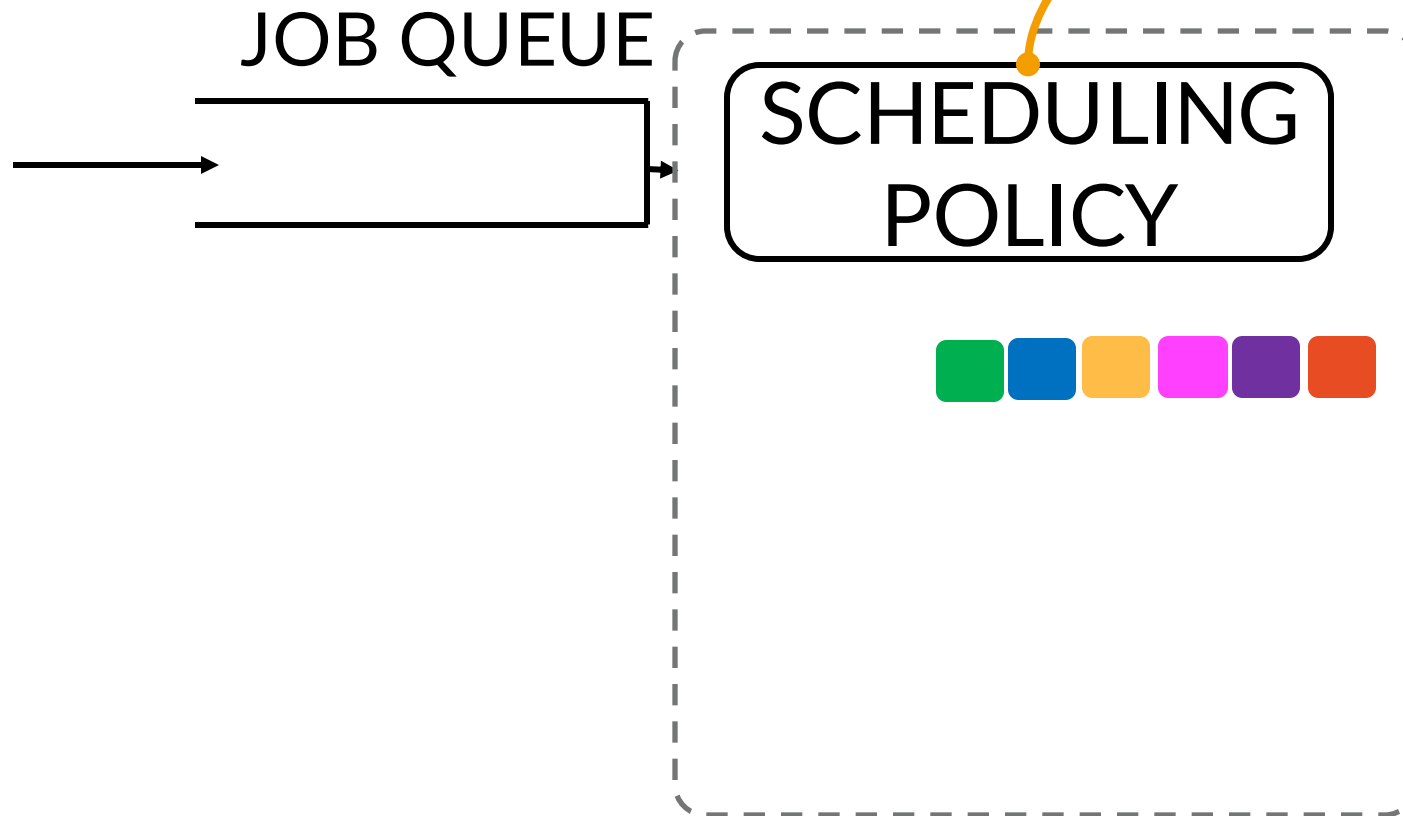
CLUSTER SCHEDULING



First In First Out (FIFO)
Least Attained Service (LAS)
Shortest Job First (SJF)



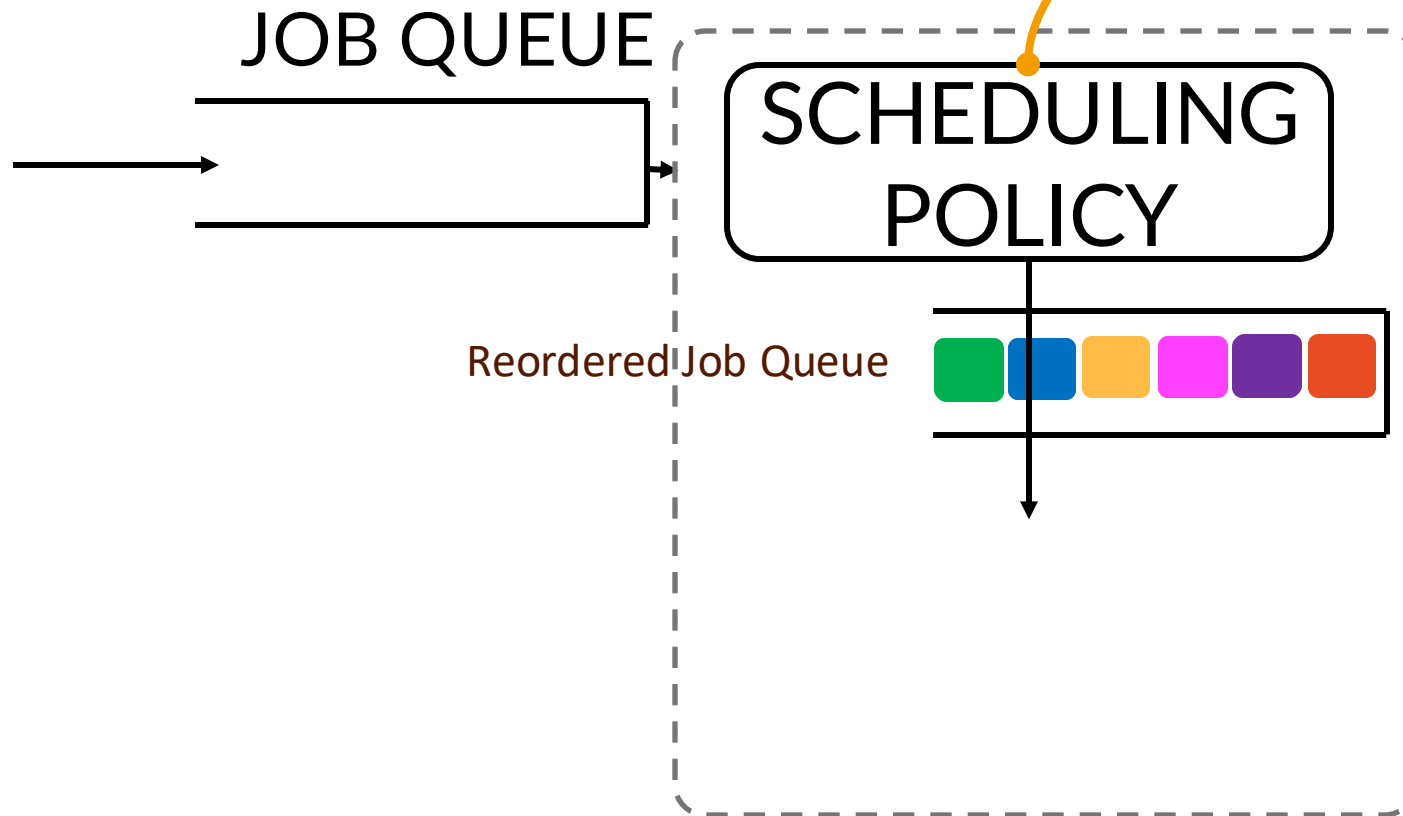
CLUSTER SCHEDULING



First In First Out (FIFO)
Least Attained Service (LAS)
Shortest Job First (SJF)



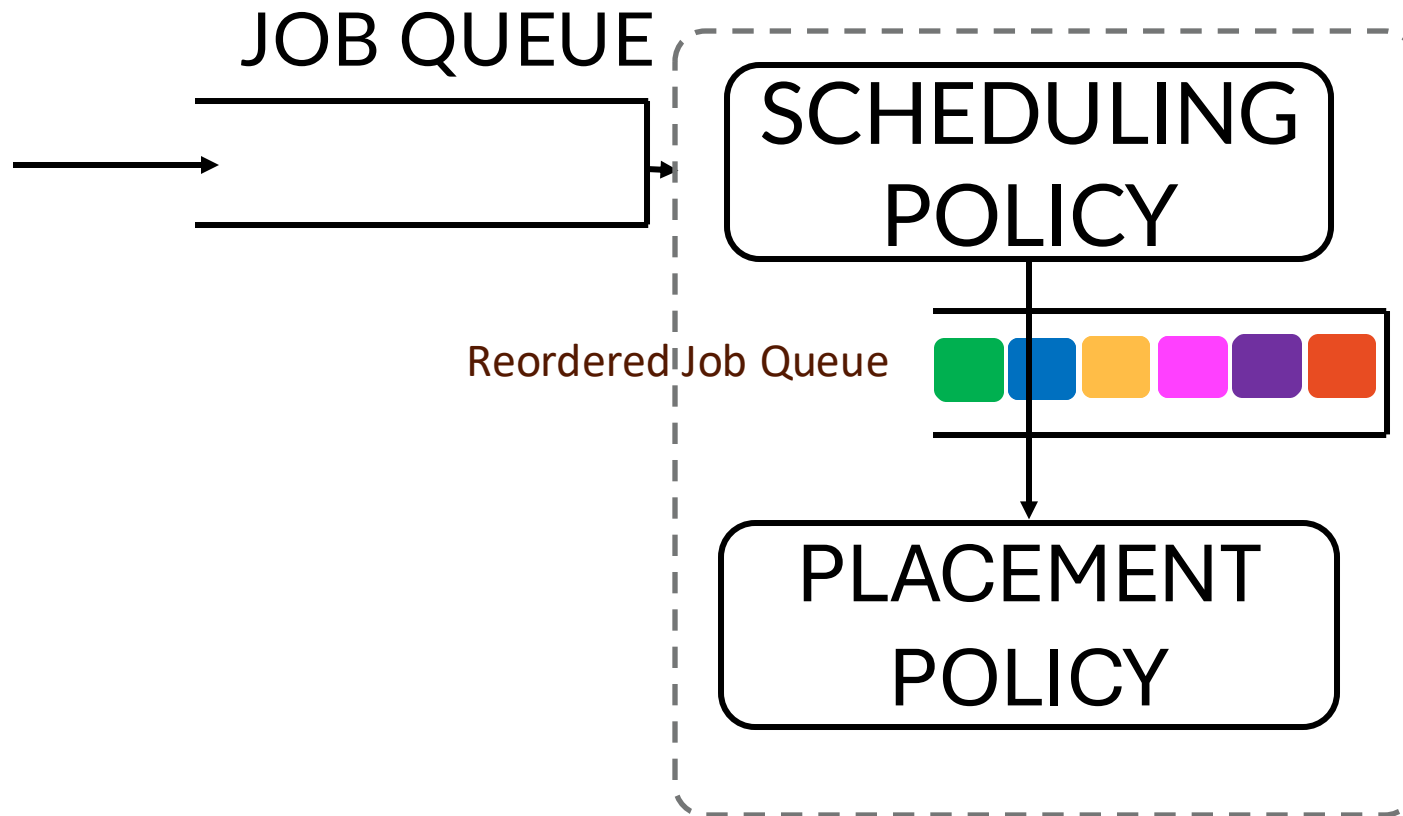
CLUSTER SCHEDULING



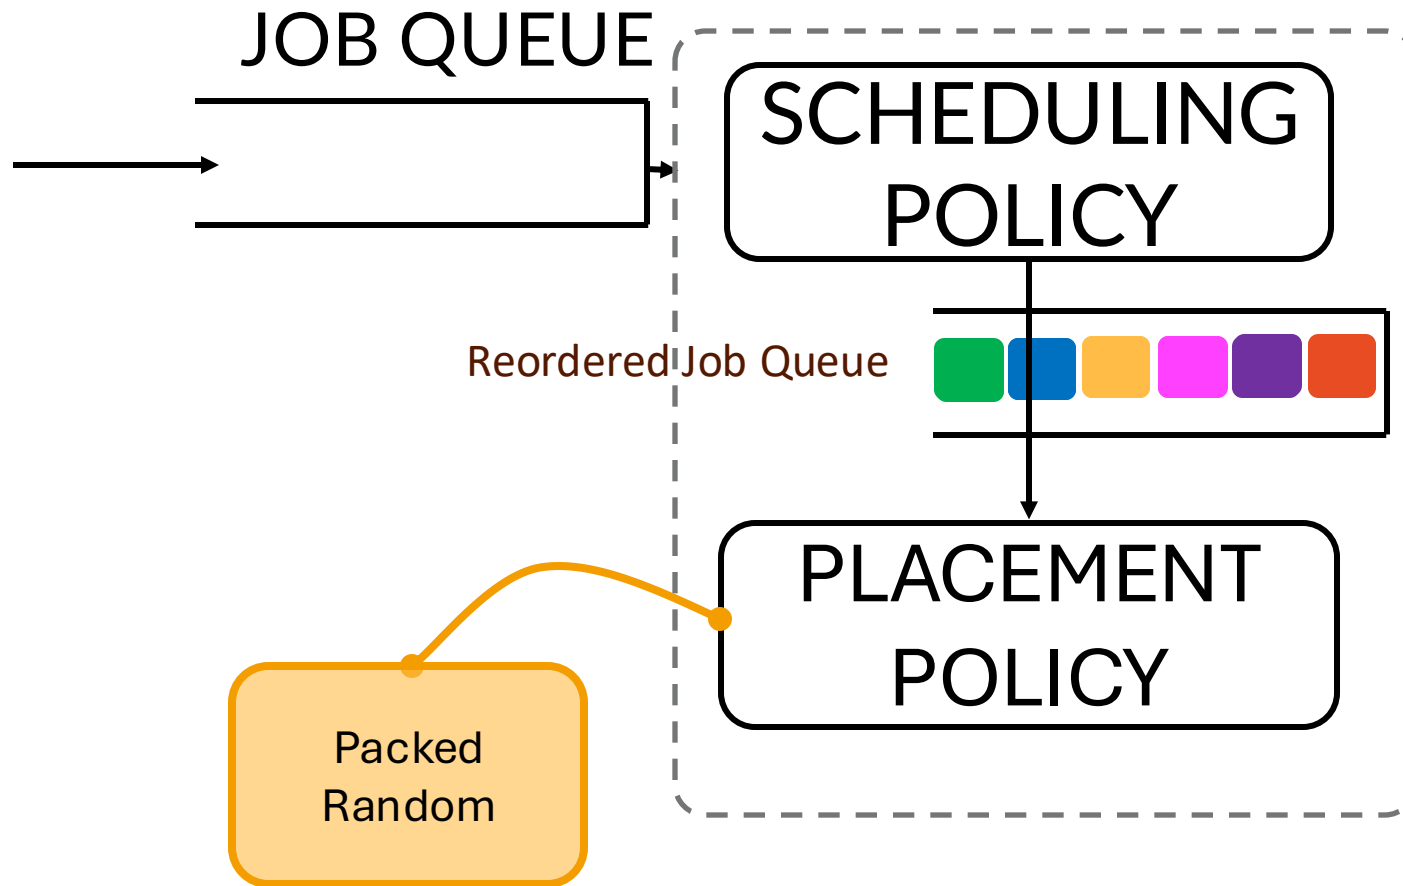
First In First Out (FIFO)
Least Attained Service (LAS)
Shortest Job First (SJF)



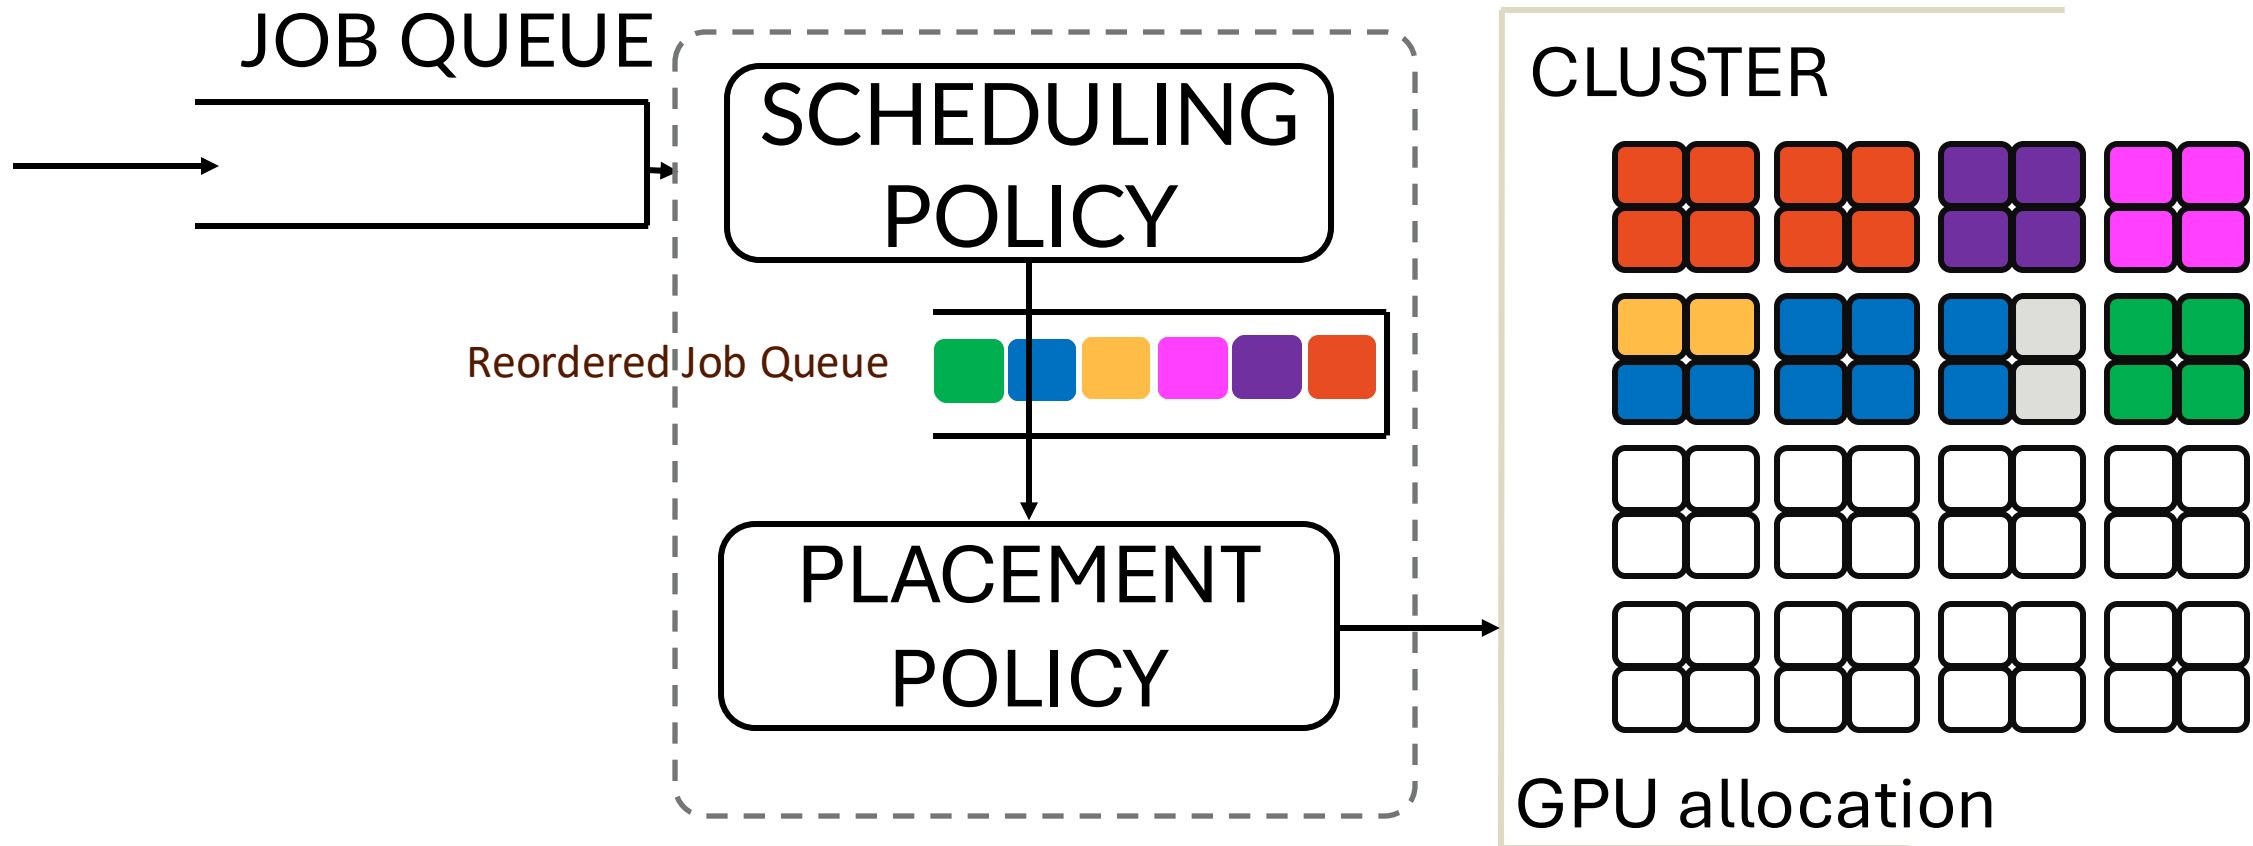
CLUSTER SCHEDULING



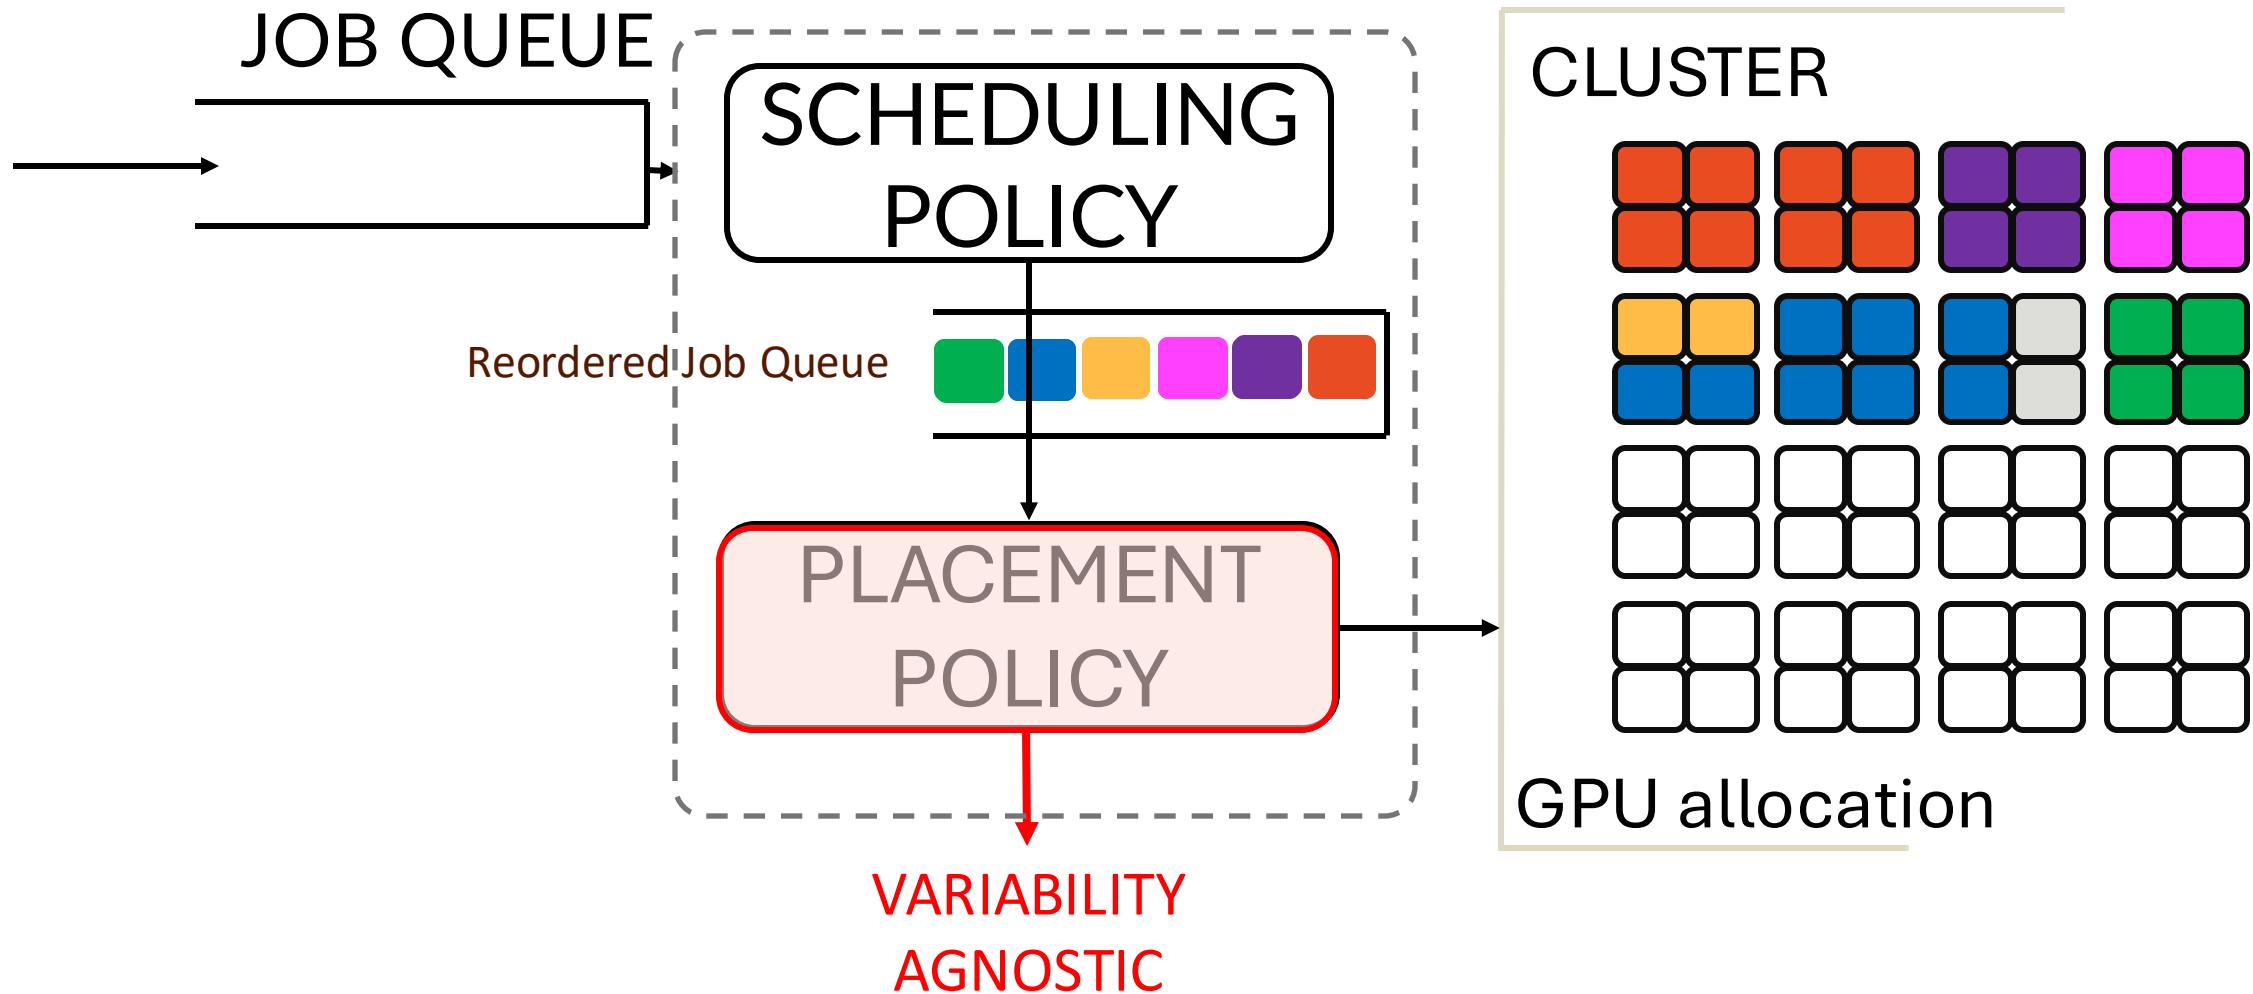
CLUSTER SCHEDULING



CLUSTER SCHEDULING

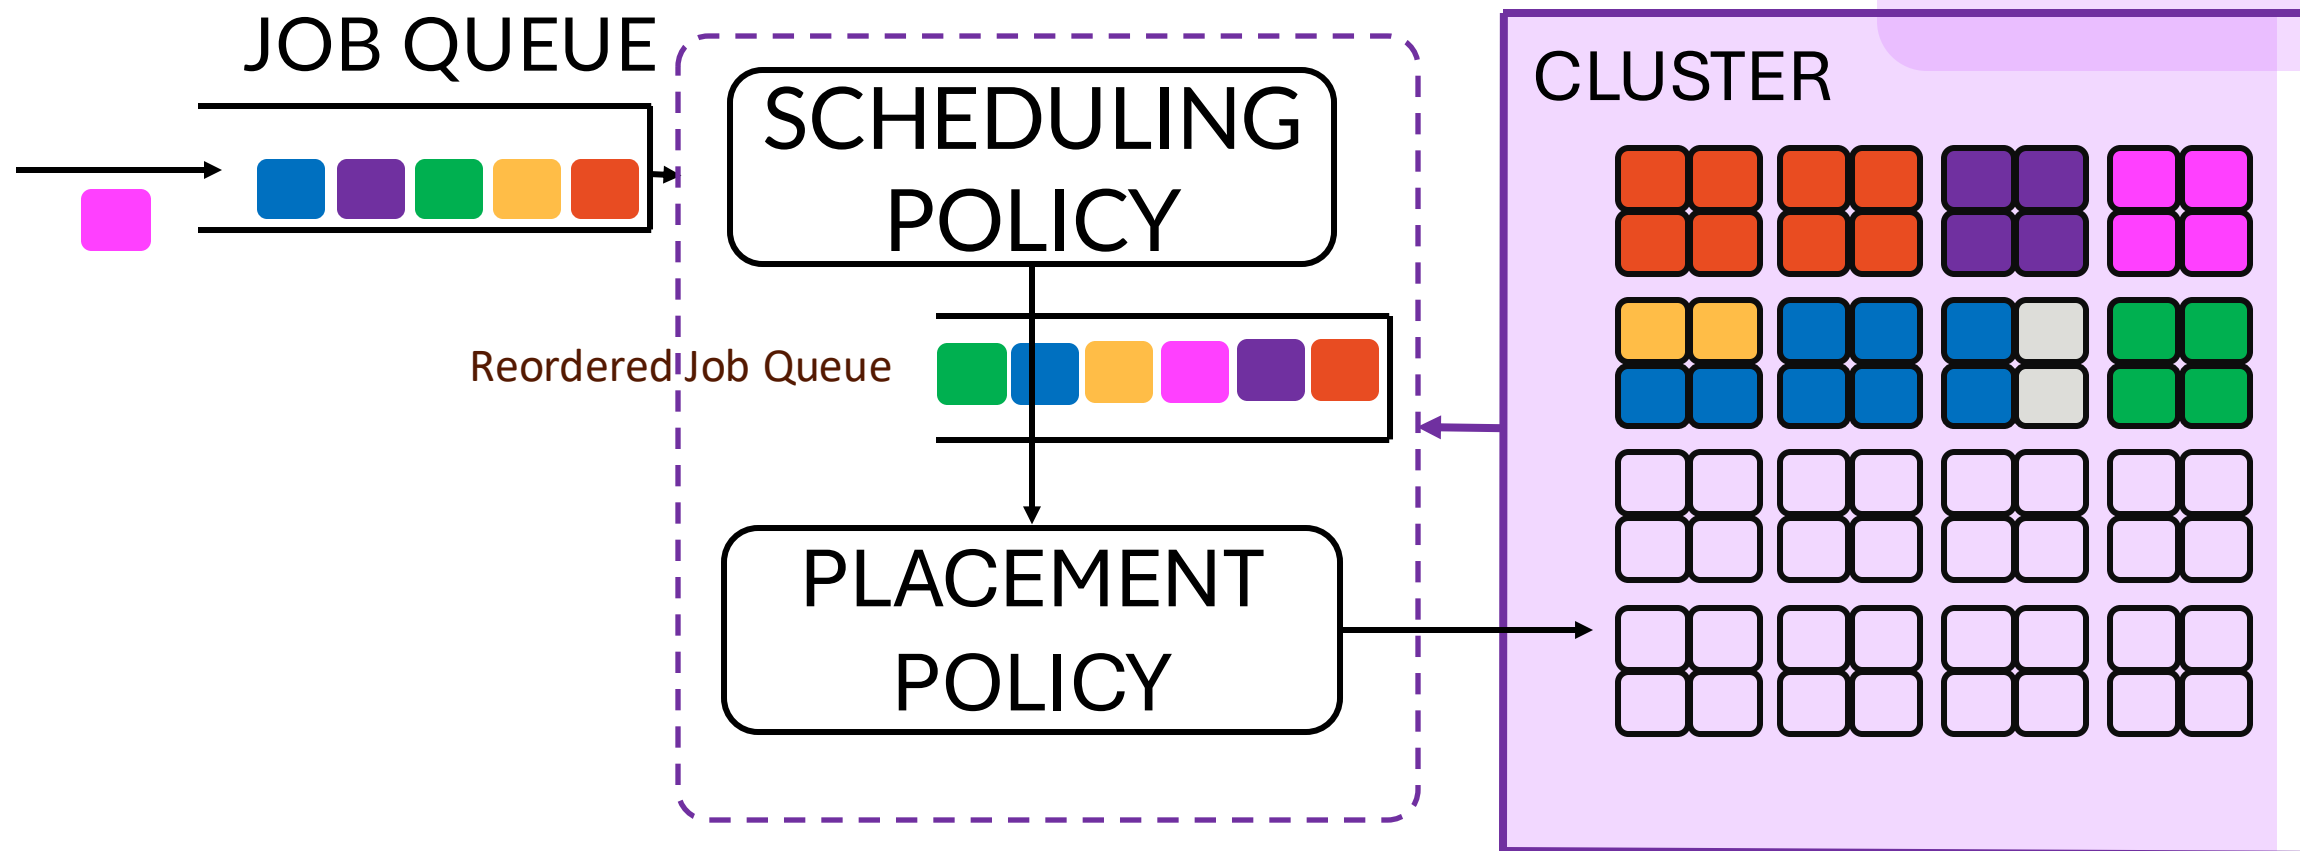


CLUSTER SCHEDULING



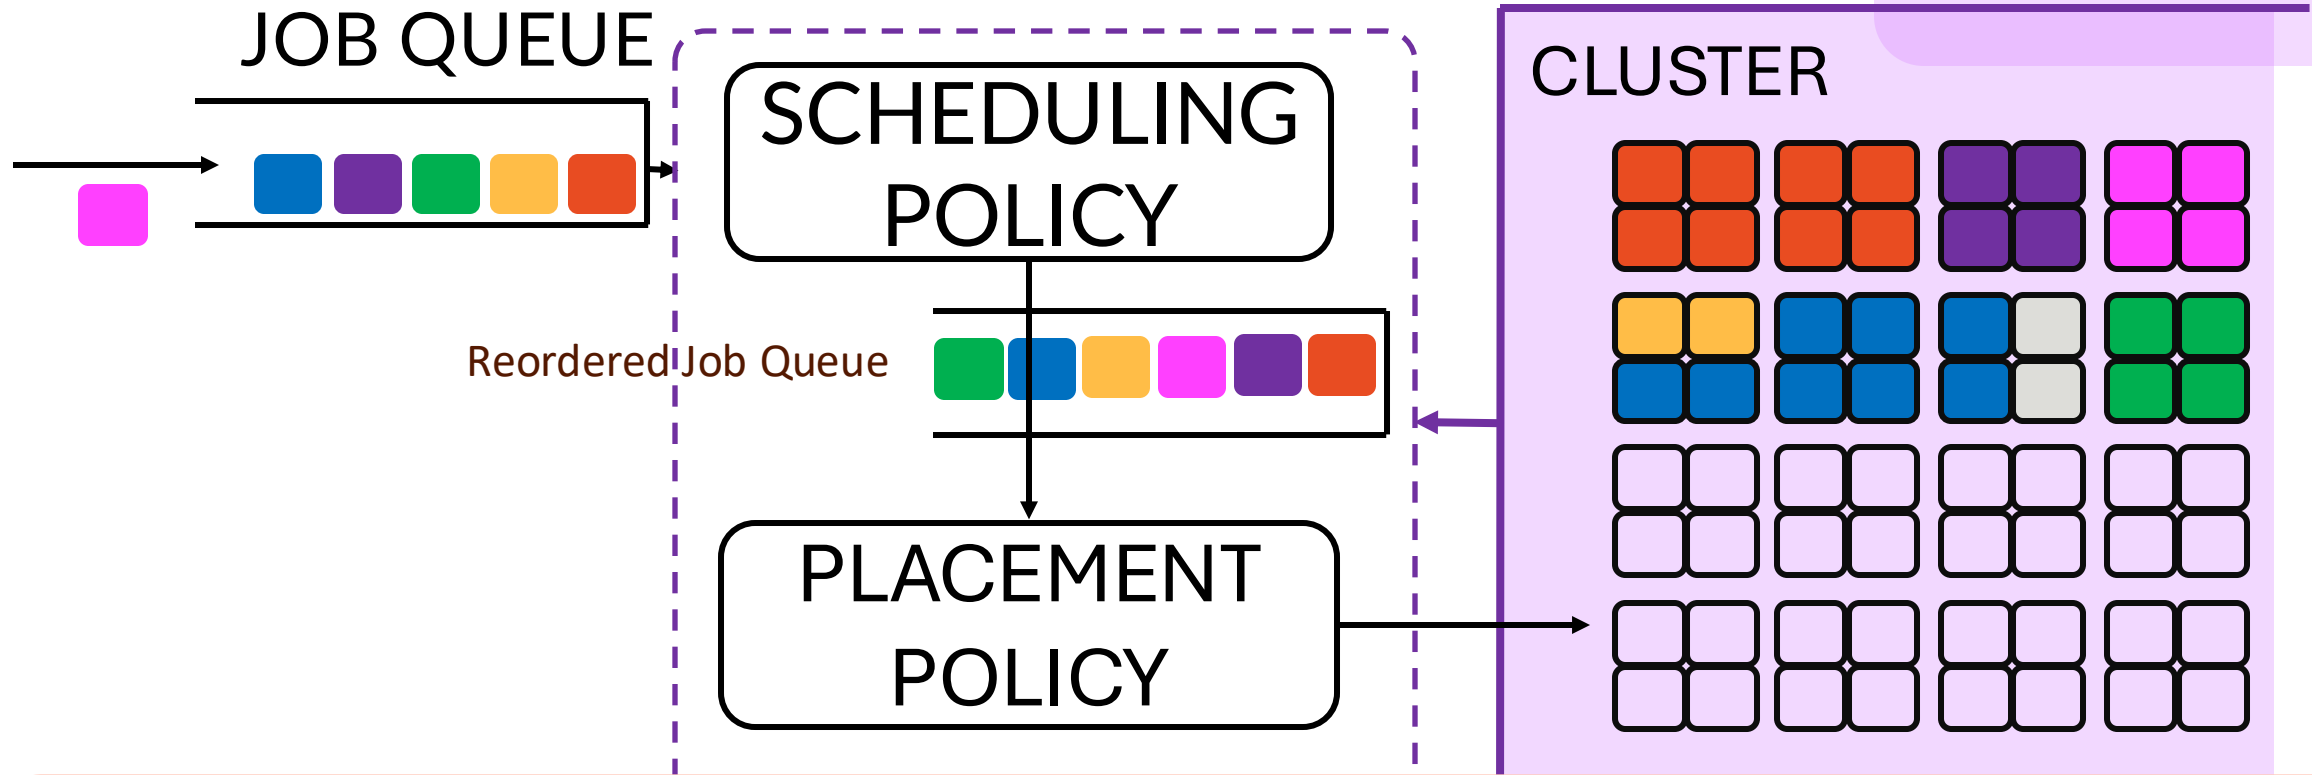
VARIABILITY-INFORMED PLACEMENT

application-specific
profiling to generate
variability data



VARIABILITY-INFORMED PLACEMENT

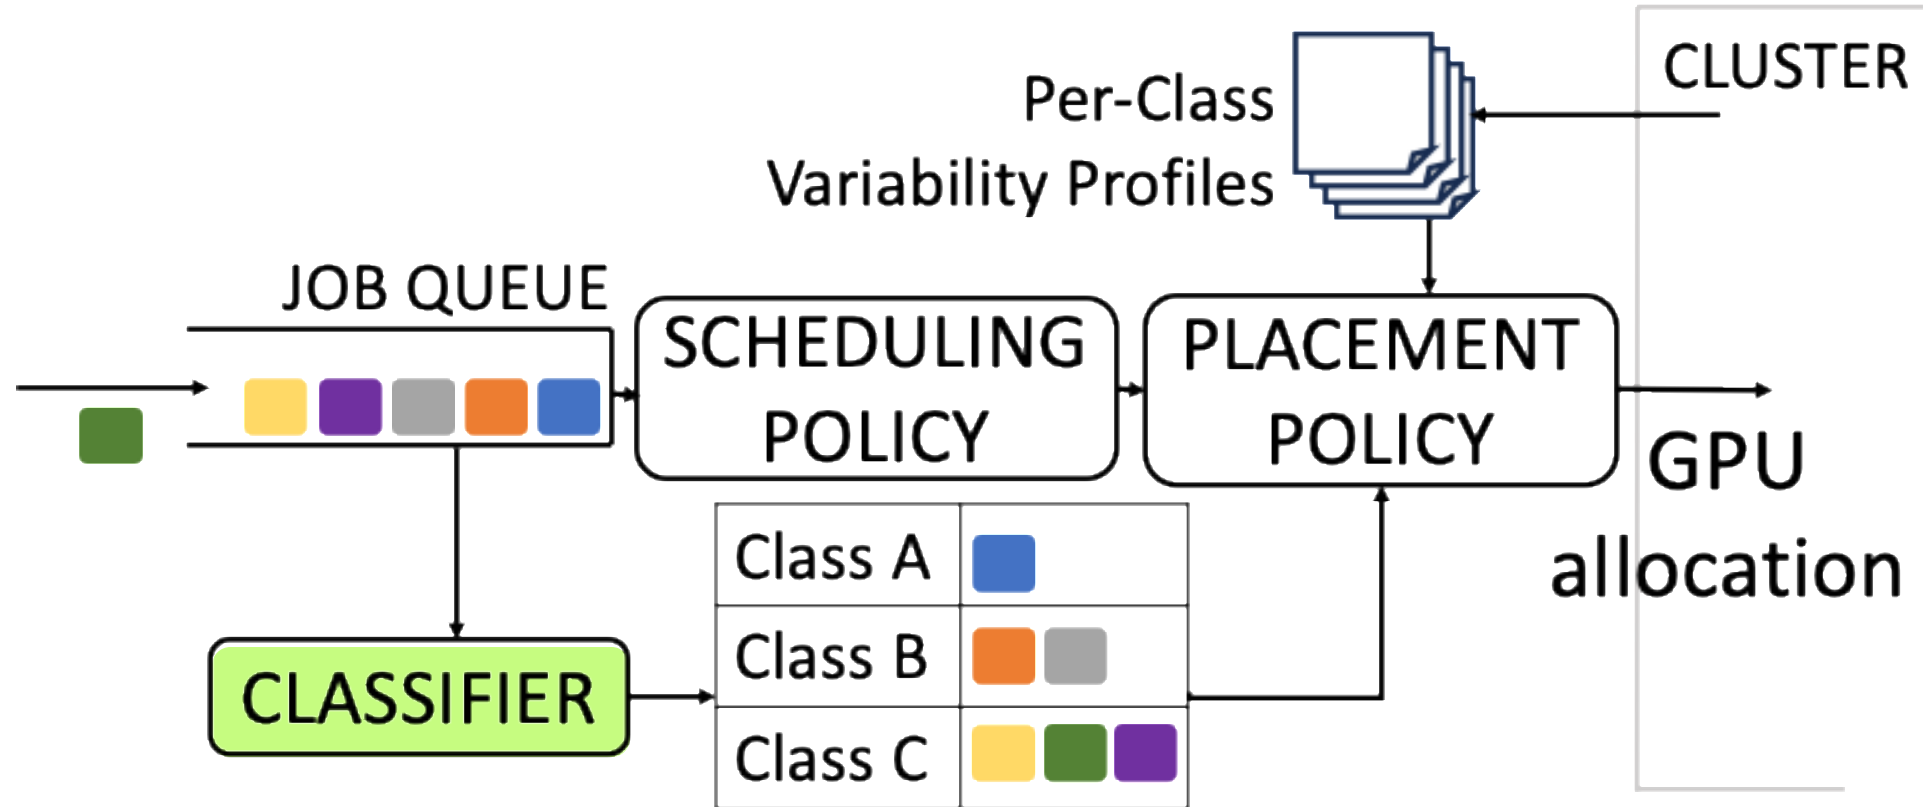
application-specific
profiling to generate
variability data



Problem: Variability is application-specific, and systems run a wide variety of workloads



VARIABILITY-INFORMED PLACEMENT

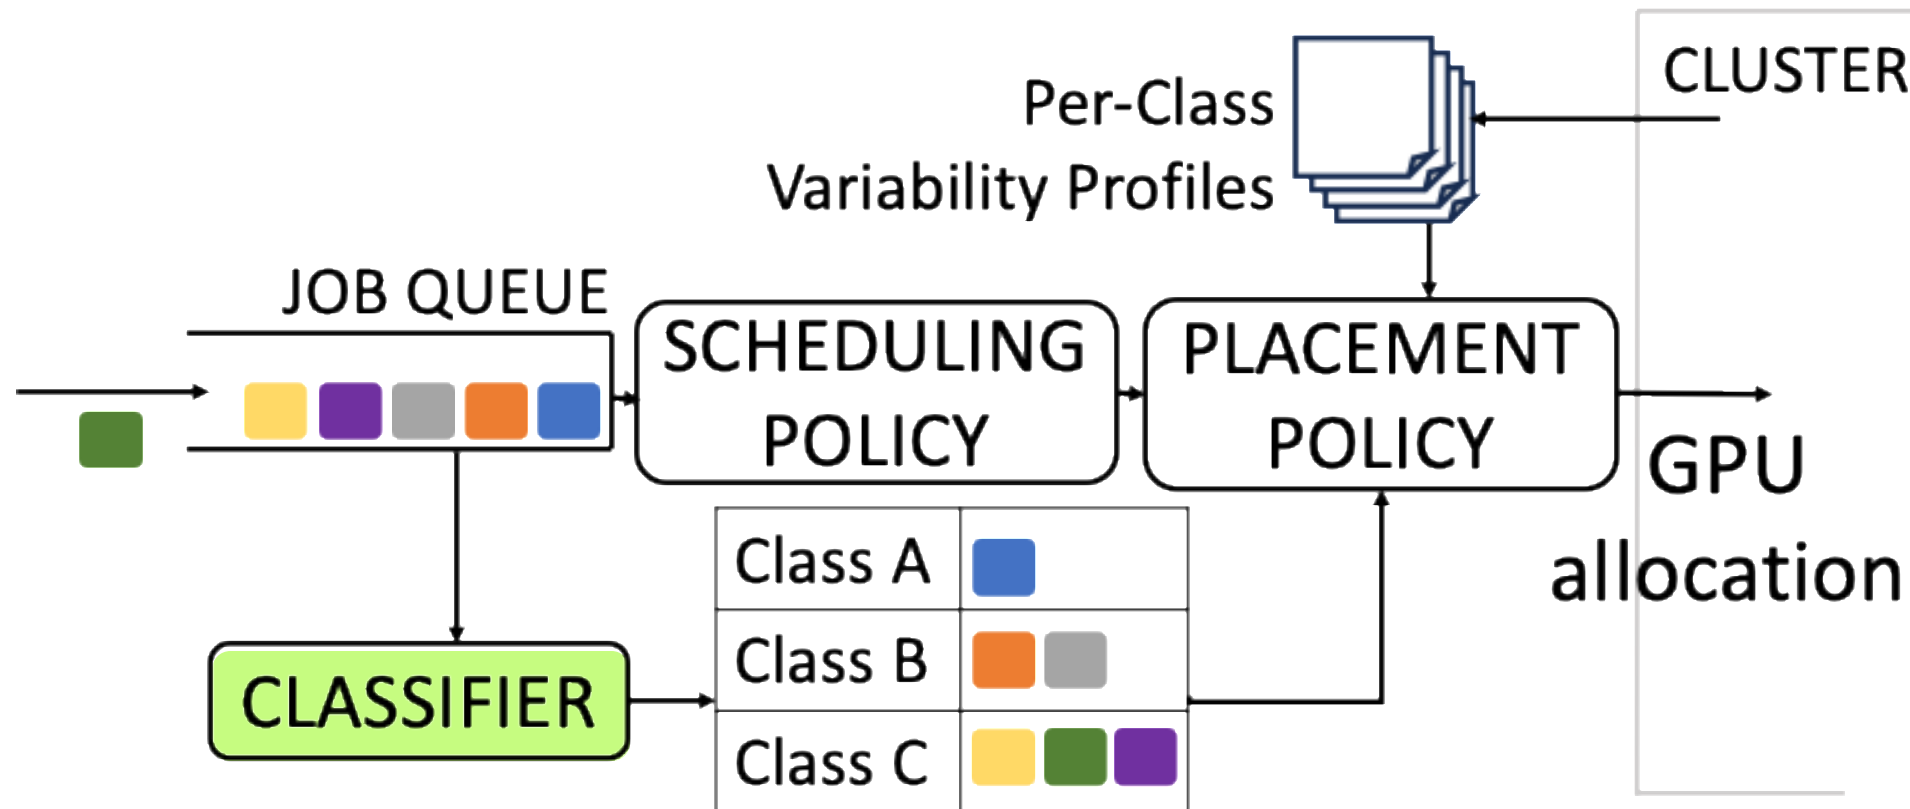


Solution: have a finite number of classes
(compute bound \leftrightarrow memory bound)



VARIABILITY-INFORMED PLACEMENT


use one variability profile per “class” as proxy for variability behavior of all applications in that class



Solution: have a finite number of classes
(compute bound \leftrightarrow memory bound)



CLASSIFICATION LAYER



- Similar to Guerreiro et al. [Parallel Computing 2019], we used **nsight compute** to measure workloads' DRAM Throughput and Peak SM Throughput.
- **2D K-Means clustering** to produce ordered classes.
- Any new application can fall into one of these classes based on its DRAM and SM throughput.



SOLUTION 1: PM-FIRST PLACEMENT



SOLUTION 1: PM-FIRST PLACEMENT

- Give variability first-order precedence when assigning GPUs to jobs.



SOLUTION 1: PM-FIRST PLACEMENT

- Give variability first-order precedence when assigning GPUs to jobs.
- We associate a **PM-score** with each GPU
 - How fast/slow the GPU is relative to median GPU in the cluster (normalized performance)

$$P_i = \frac{t_i}{t_{median}}$$

- PM-score is class-specific.



SOLUTION 1: PM-FIRST PLACEMENT

- Give variability first-order precedence when assigning GPUs to jobs.
- We associate a **PM-score** with each GPU
 - How fast/slow the GPU is relative to median GPU in the cluster (normalized performance)

$$P_i = \frac{t_i}{t_{median}}$$

- PM-score is class-specific.
- We also associate a **placement priority** for each job




SOLUTION 1: PM-FIRST PLACEMENT

- Give variability first-order precedence when assigning GPUs to jobs.
- We associate a **PM-score** with each GPU
 - How fast/slow the GPU is relative to median GPU in the cluster (normalized performance)

$$P_i = \frac{t_i}{t_{median}}$$

- PM-score is class-specific.
- We also associate a **placement priority** for each job




SOLUTION 1: PM-FIRST PLACEMENT

- Give variability first-order precedence when assigning GPUs to jobs.
- We associate a **PM-score** with each GPU
 - How fast/slow the GPU is relative to median GPU in the cluster (normalized performance)

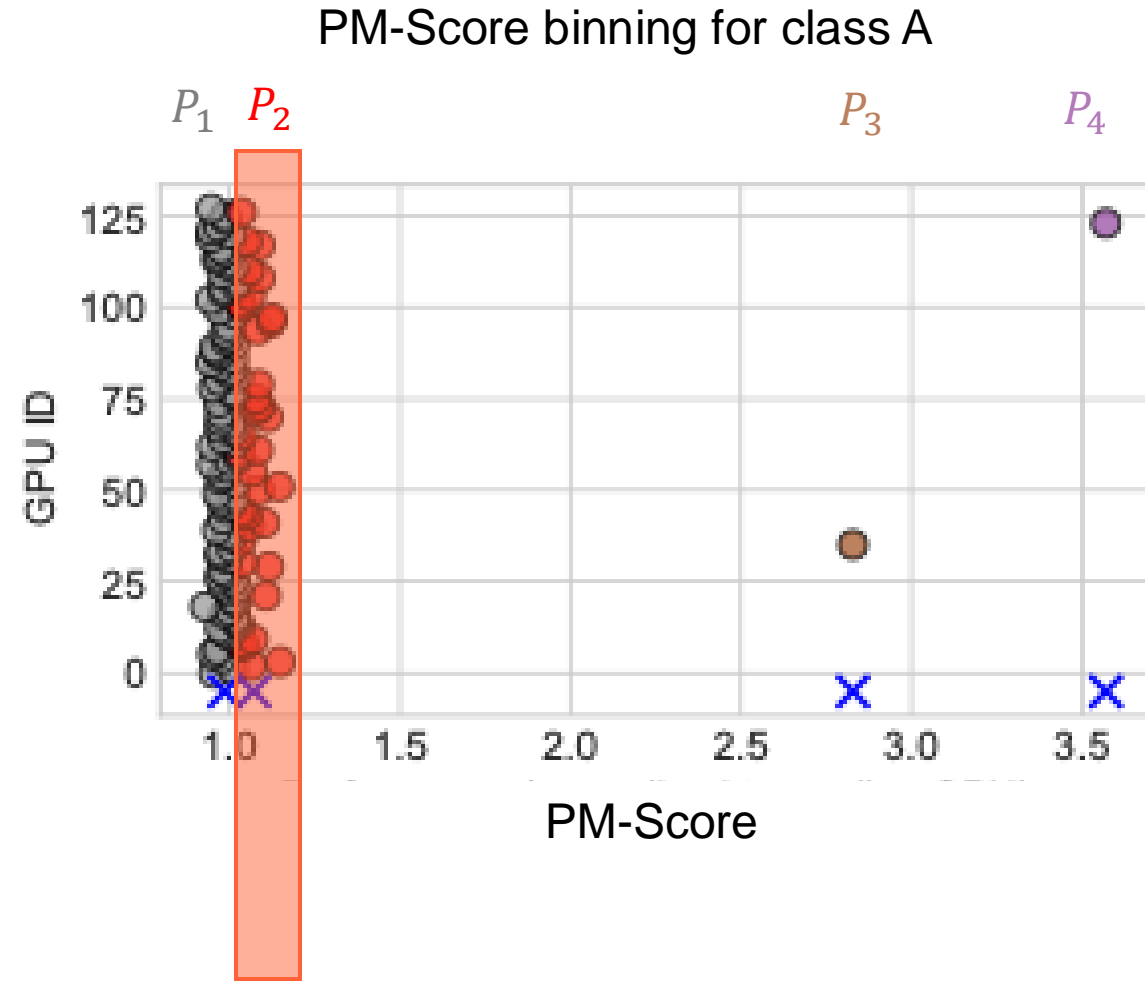
$$P_i = \frac{t_i}{t_{median}}$$

- PM-score is class-specific.
- We also associate a **placement priority** for each job



SOLUTION 1: PM-FIRST PLACEMENT

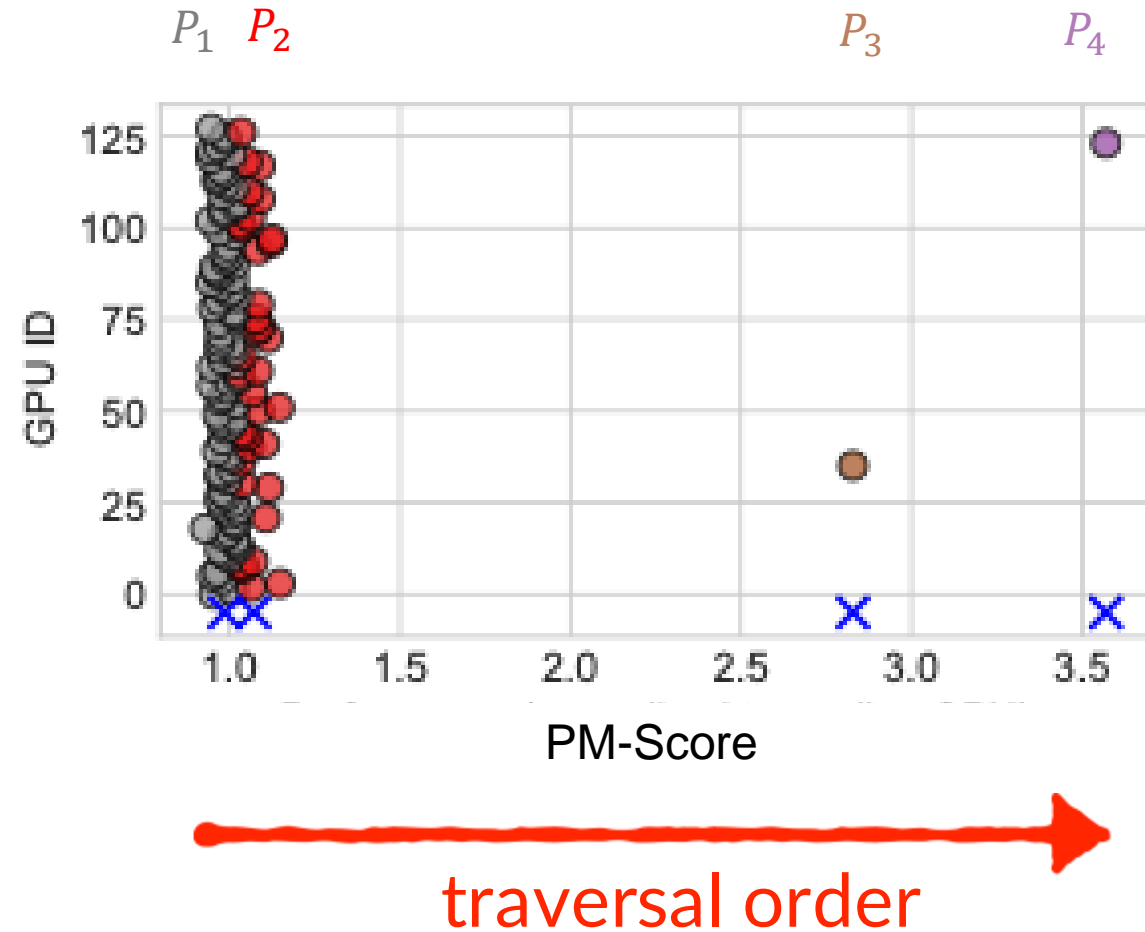
- For each class, we make **bins** of GPUs with similar **PM-Scores**.
- Why?
 - Fine-grained variability for large-scale systems is expensive
 - Memory-bound applications have little variability – one large pool of GPUs for them.



all GPUs of the red bin have similar variability



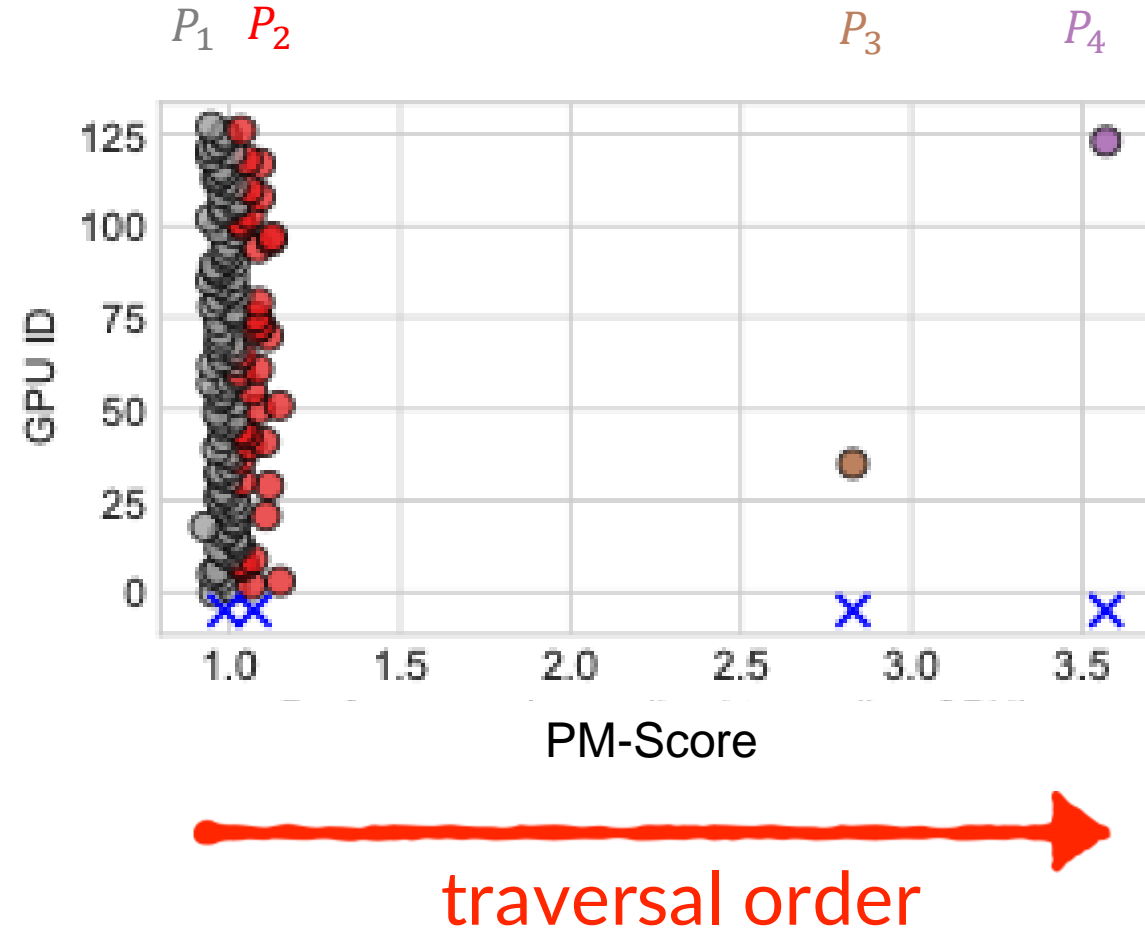
PM-FIRST PLACEMENT POLICY



PM-FIRST PLACEMENT POLICY

In each scheduling round,

- Get job queue from scheduling policy
- Mark queue at cluster size, and assign placement priority




PM-FIRST PLACEMENT POLICY

In each scheduling round,

- Get job queue from scheduling policy
- Mark queue at cluster size, and assign placement priority

for each job j in order of placement priority:

- get job **class** and GPU demand N
- get **freelist** of GPUs available and **sort** by their **PM-Score** value (for job **class**)
- allocate first N GPUs from this sorted list.
- Mark GPUs in use (remove from **freelist**)



WHAT ABOUT LOCALITY?

- PM-First ignores communication overheads due to ineffective packing
- Some jobs are more sensitive to locality of GPUs than variability



WHAT ABOUT LOCALITY?

- PM-First ignores communication overheads due to ineffective packing
- Some jobs are more sensitive to locality of GPUs than variability
- Assume a simplified locality model:
 - Multi-GPU job incurs penalty L_{across} if it's allocation spills across nodes
 - No performance degradation if allocation is within a node
- Modified iteration time for a job running with set of GPUs G spread across > 1 node:

$$t_{iter} = L_{across} \times \max_{g \in G}(P_g) \times t_{iter}^{orig}$$

where P_g is the PM-Score of GPU g



WHAT ABOUT LOCALITY?

- PM-First ignores communication overheads due to ineffective packing
- Some jobs are more sensitive to locality of GPUs than variability
- Assume a simplified locality model:
 - Multi-GPU job incurs penalty L_{across} if it's allocation spills across nodes
 - No performance degradation if allocation is within a node
- Modified iteration time for a job running with set of GPUs G spread across > 1 node:

$$t_{iter} = L_{across} \times \max_{g \in G}(P_g) \times t_{iter}^{orig}$$

where P_g is the PM-Score of GPU g

$$\downarrow t_{iter} = \downarrow (L_{across} \times V) \times t_{iter}^{orig}$$



LV-PRODUCT

$$\min LV\text{-product} = \min (L_{across} \times V)$$

↓
Slowdown due to
locality

↓
Slowdown due to
variability

Lower V values are better

Example ClassA	$V_1 = 0.89$	$V_2 = 0.94$	$V_3 = 1.06$	$V_4 = 2.55$
Within Node $L_{within} = 1$	0.89	0.94	1.06	2.55
Across node $L_{across} = 1.5$	1.35	1.41	1.59	3.82



LV-PRODUCT

$$\min LV\text{-product} = \min (L_{across} \times V)$$

↓
Slowdown due to
locality

↓
Slowdown due to
variability

Lower V values are better

Example ClassA	$V_1 = 0.89$	$V_2 = 0.94$	$V_3 = 1.06$	$V_4 = 2.55$
Within Node $L_{within} = 1$	0.89	0.94	1.06	2.55
Across node $L_{across} = 1.5$	1.35	1.41	1.59	3.82

PAL traversal order



LV-PRODUCT

$$\min LV\text{-product} = \min (L_{across} \times V)$$

↓
Slowdown due to
locality

↓
Slowdown due to
variability

Example ClassA	$V_1 = 0.89$	$V_2 = 0.94$	$V_3 = 1.06$	$V_4 = 2.55$
Within Node $L_{within} = 1$	0.89	0.94	1.06	2.55
Across node $L_{across} = 1.5$	1.35	1.41	1.59	3.82

PM-First traversal order



OUTLINE

- Introduction
- Cluster Scheduling
- Design
- **Methodology**
- Evaluation
- Conclusion



SYSTEM

- **Blox (open-source modular toolkit)**
- **Physical Cluster**



TACC Frontera

360 GPUs

mineral oil cooled



NVIDIA Quadro RTX5000 GPUs

4x GPUs per node

16GB memory per GPU

- **Testbed**
 - 16 node (64 GPUs)



SYSTEM

- Blox (open-source modular toolkit)

- Physical Cluster



TACC Frontera

360 GPUs

mineral oil cooled



NVIDIA Quadro RTX5000 GPUs

4x GPUs per node

16GB memory per GPU

- Testbed

- 16 node (64 GPUs)

Experiment(s)	Cluster Size	Workload Trace	Paper Section
Cluster Testbed Evaluation	64	Sia-Cluster*	V-A
Baseline Simulation Varying Locality Penalty	64	Sia-Philly*	V-B
Varying Job Load Varying Scheduling Policy	256	Synergy†	V-C

* [Subramanya et al. SOSP'23]

† [Mohan et al. OSDI'22]




BASELINE PLACEMENT POLICIES



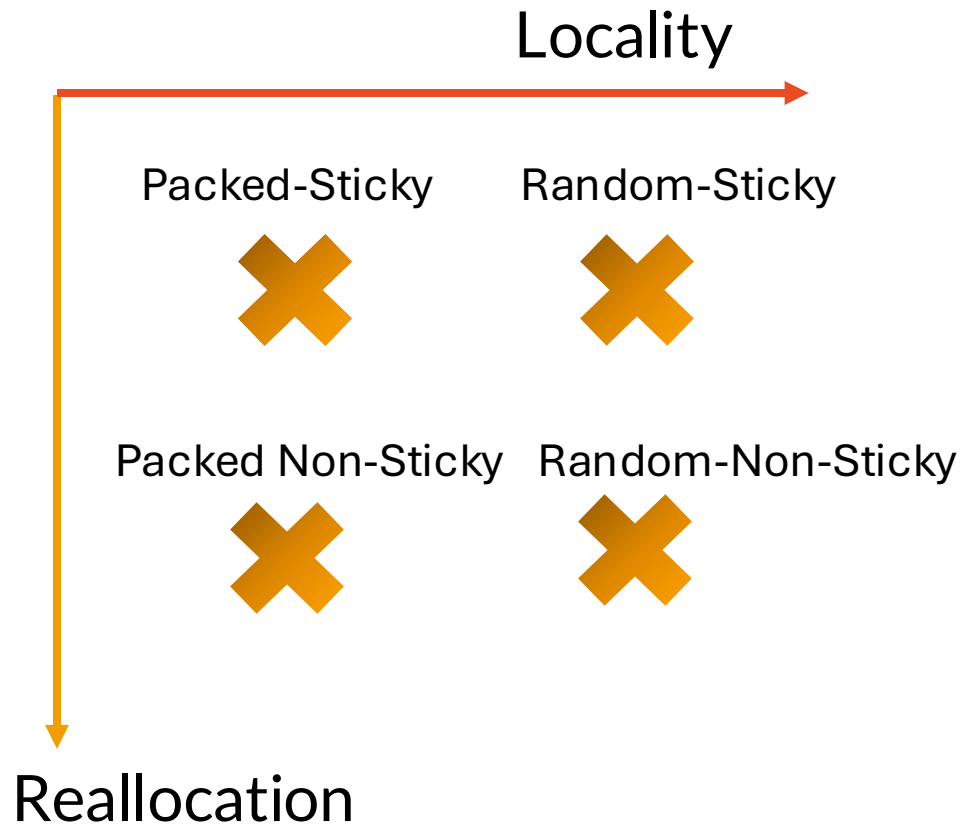
BASELINE PLACEMENT POLICIES




BASELINE PLACEMENT POLICIES



METHODOLOGY: PLACEMENT POLICIES



METHODOLOGY: PLACEMENT POLICIES

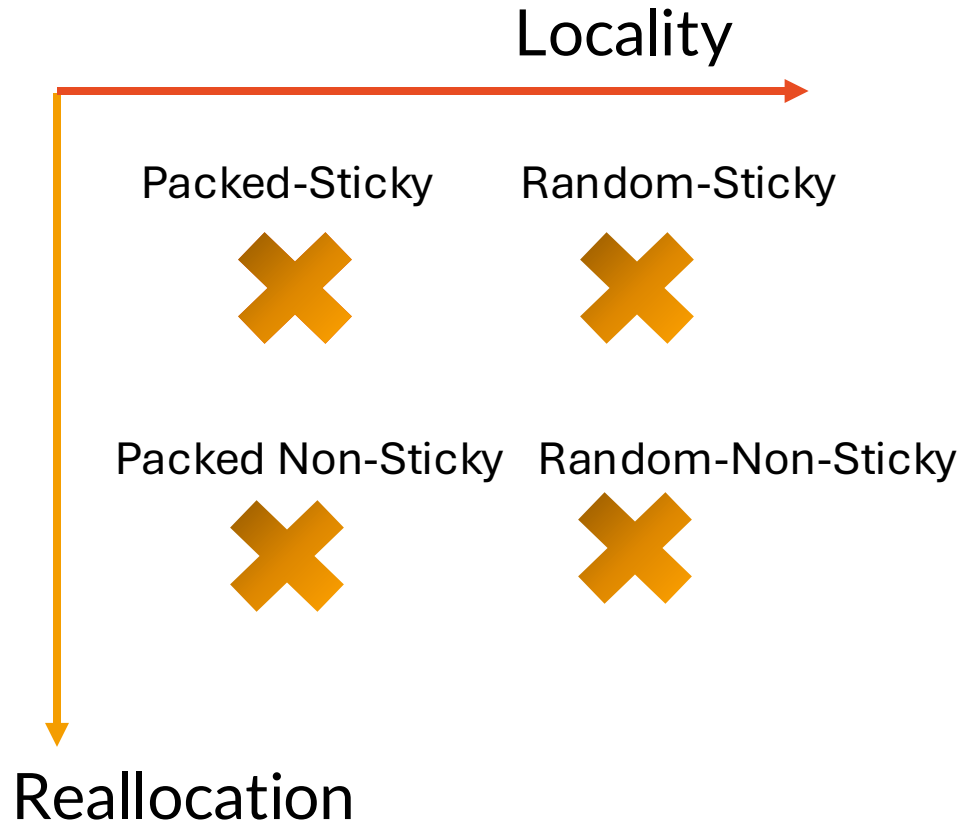


Baseline placement policies

- **Tiresias** [Gu et al. NSDI'19]
 - Packed Sticky
- **Gandiva** [Xia et al. OSDI'18]
 - Packed Non-sticky
- **Random-Sticky** HotGauge
- **Random-Non-Sticky** [Hankin et al. IISWC'21]



METHODOLOGY: PLACEMENT POLICIES



Baseline placement policies

- **Tiresias** [Gu et al. NSDI'19]
 - Packed Sticky
 - **Gandiva** [Xia et al. OSDI'18]
 - Packed Non-sticky
 - **Random-Sticky** HotGauge
 - **Random-Non-Sticky** [Hankin et al. IISWC'21]
-
- **PM-First and PAL**
 - Both non-sticky




APPLICATIONS: PHYSICAL CLUSTER EXPERIMENTS

Task	Model	Dataset	Batch Size	Class
Image	PointNet	ShapeNet	32	C
Image	vgg19	ImageNet2012	32	A
Vision	DCGAN	LSUN	128	A
Language	BERT	WikiText	64	B
Image	ResNet-50	Imagenet2012	32	A
Language	GPT2	Wikitext	128	B




PM-SCORES: VARIABILITY PROFILES




$$P_i = \frac{t_i}{t_{median}}$$



LOCALITY PENALTY




8 GPU ResNet-50 job
Batch Size 128



4 GPU ResNet-50 job
Batch Size 64




LOCALITY PENALTY



8 GPU ResNet-50 job
Batch Size 128



$t_{iter,avg}$



4 GPU ResNet-50 job
Batch Size 64



$t_{iter,avg}$

$$\frac{\quad}{\quad} = L_{across}$$





OUTLINE


- Introduction
- Cluster Scheduling
- Design
- Methodology
- **Evaluation**
- Conclusion





PHYSICAL CLUSTER EXPERIMENT: 64 GPUs



PHYSICAL CLUSTER EXPERIMENT: 64 GPUs




PHYSICAL CLUSTER EXPERIMENT: 64 GPUs



cluster-to-sim diff
11-14%



PHYSICAL CLUSTER EXPERIMENT: 64 GPUs



First-order trends are accurate and correlate well in simulation




SIA-PHILLY SIMULATIONS

Scheduler: **FIFO**

of GPUs: **64**

Job trace: **Sia-Philly trace (160 jobs, requesting up to 48 GPUs)**




SIA-PHILLY SIMULATIONS

Scheduler: **FIFO**

of GPUs: **64**

Job trace: **Sia-Philly trace (160 jobs, requesting up to 48 GPUs)**



PM-First improves geomean JCT by 40%, PAL further improves JCTs by considering locality in addition to variability




SIA-PHILLY SIMULATIONS

Scheduler: **FIFO**

of GPUs: **64**




Job trace: **Sia-Philly trace (160 jobs, requesting up to 48 GPUs)**





PM-First improves geomean JCT by 40%, PAL further improves JCTs by considering locality in addition to variability





SIA-PHILLY SIMULATIONS



SIA-PHILLY SIMULATIONS



SIA-PHILLY SIMULATIONS




Large jobs hog cluster resources, causing cascading wait time impact!

Jobs finish faster with PAL and subsequent wait times consequently reduce

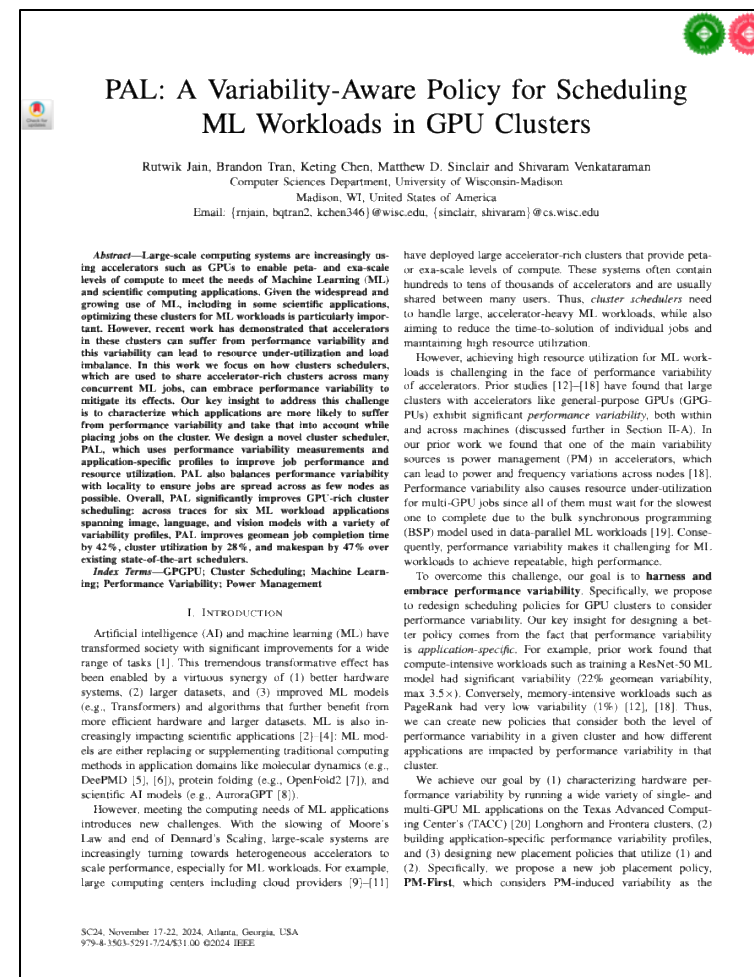


Varying Locality Penalty



job arrival rate, cluster utilization, placement overheads...

- Varying cluster size (64 vs 256) – **Section V-B vs V-C**
- Varying locality penalty – **Section V-B (1) and V-C (2)**
- Varying job load (jobs/hr) – **Section V-C (1)**
 - Average JCT
 - Utilization
- Different proportions of single GPU jobs - **Section V-B vs V-C**
- Varying scheduling policy - **Section V-C (1)**
 - FIFO
 - SRTF
 - LAS
- Placement policy overheads - **Section V-C (2)**



CONCLUSION



CONCLUSION



Variability can affect cluster performance, utilization and load balancing on GPU clusters. Likely to get worse as **ML algorithms** and **cluster size** scales up, while **transistors** shrink.



CONCLUSION



Variability can affect cluster performance, utilization and load balancing on GPU clusters. Likely to get worse as **ML algorithms** and **cluster size** scales up, while **transistors** shrink.



Key Insight: Schedulers must **embrace** and harness this variability.



CONCLUSION



Variability can affect cluster performance, utilization and load balancing on GPU clusters. Likely to get worse as **ML algorithms** and **cluster size** scales up, while **transistors** shrink.




Key Insight: Schedulers must **embrace** and harness this variability.



We proposed **PAL** which uses application-specific variability to optimize job placement. Across a mix of ML workloads, PAL improves on SOTA ML cluster schedulers across metrics.



CONCLUSION



Variability can affect cluster performance, utilization and load balancing on GPU clusters. Likely to get worse as **ML algorithms** and **cluster size** scales up, while **transistors** shrink.



Key Insight: Schedulers must **embrace** and harness this variability.



We proposed **PAL** which uses application-specific variability to optimize job placement. Across a mix of ML workloads, PAL improves on SOTA ML cluster schedulers across metrics.



We are working to extend this to HPC and HPC+ML workloads. Looking at HW-SW codesign with variability as a first-order constraint.



SUMMARY



Variability can affect cluster performance, utilization and load balancing on GPU clusters
Likely to get worse as ML algorithms and cluster size scales up, while transistors shrink



Key Insight: Schedulers must embrace and harness this variability



We proposed **PAL** which uses application-specific variability to optimize job placement
Across a mix of ML workloads, PAL improves on SOTA ML cluster schedulers across metrics.



We are working to extend this to HPC and HPC+ML workloads.
Looking at HW-SW codesign with variability as a first-order constraint.

Contact



rnjain@wisc.edu



<https://pages.cs.wisc.edu/~rnjain/>

Artifact




THANK YOU
Questions?



Placement Policy Overheads

- Scheduling round duration is 300s (5 mins)
- PAL
 - worst-case 4 seconds
 - median of 2.8 seconds
- PM-First
 - worst-case 2 seconds




Blox Cluster Setup

Scheduler Node: CPU Host

- gRPC Server that registers all GPU UUIDs.
- Another CPU threads sends job requests at specified arrival times.
- Scheduler thread runs placement policy algorithm to determine set of GPUs to run on.
- Server dispatches job by sending a gRPC packet to the clients running on relevant node(s).


GPU Cluster



Execution Node

- Host on each execution node acts as gRPC client
- At startup, client registers its GPUs with the scheduler node and receives acknowledgement.
- It then waits for job request packets from server
- Whenever it receives a job request, it runs corresponding job command on specified list of GPUs.





How do we define variability?

Interquartile Range

$$\text{IQR} = Q_3 - Q_2$$


Whiskers and Range

$$W_2 = Q_3 + 1.5 \times \text{IQR}$$

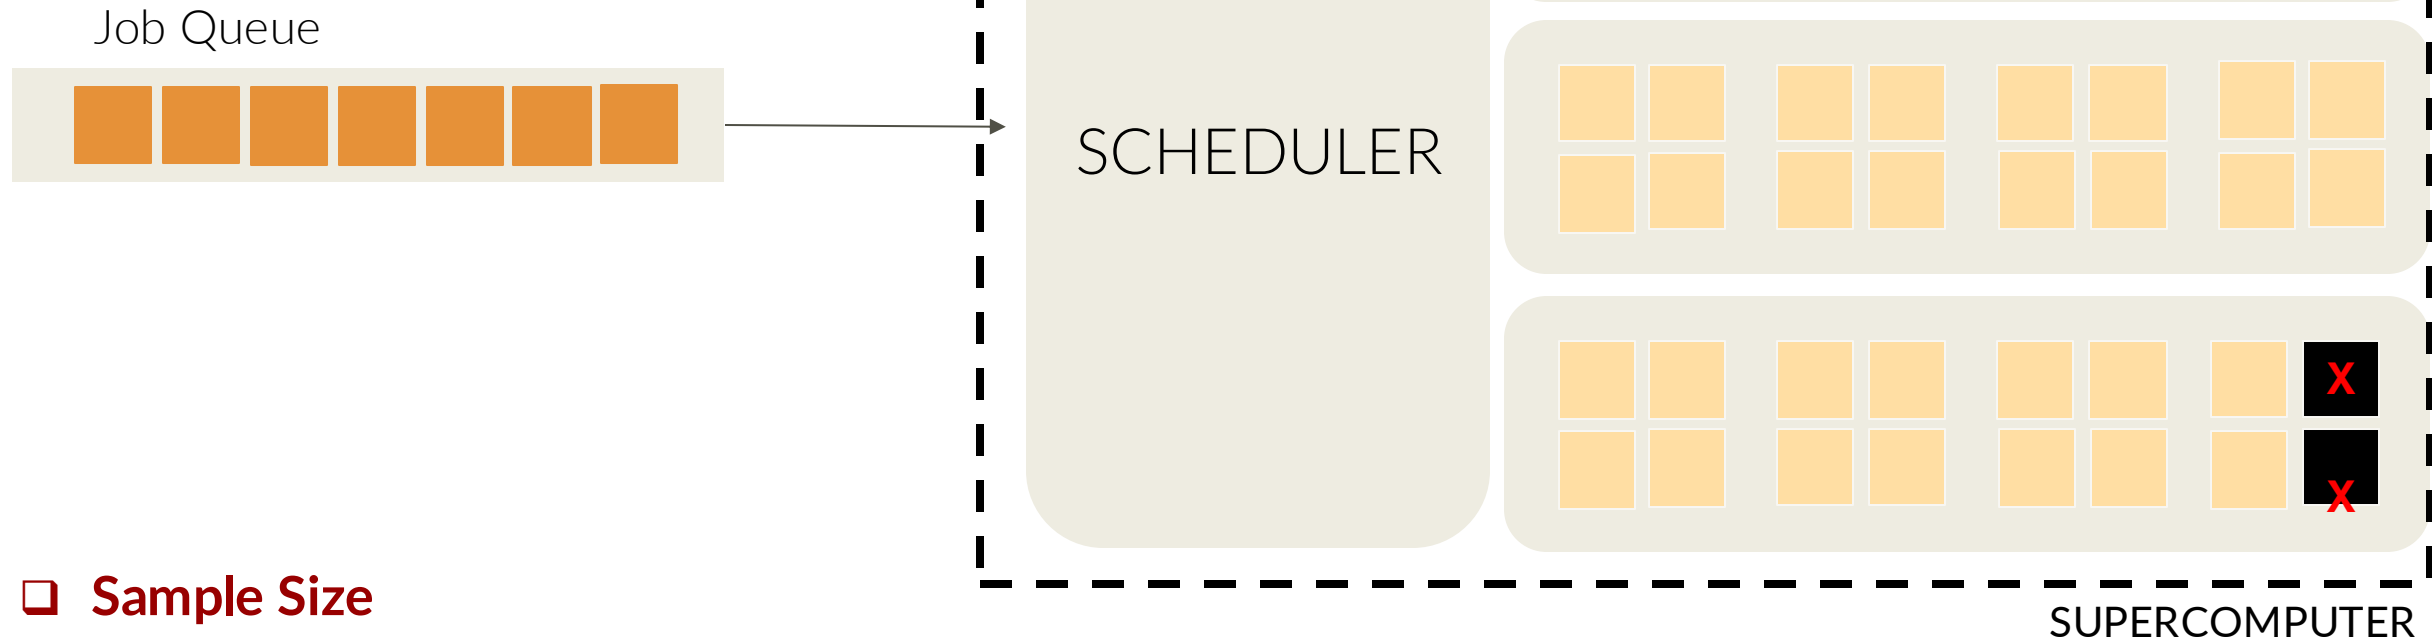
$$W_1 = Q_1 - 1.5 \times \text{IQR}$$

$$\text{Range}_{metric} = W_2 - W_1$$

$$\text{Variability}_{metric} = \frac{\text{Range}_{metric}}{\text{Median}_{metric}}$$



Methodology



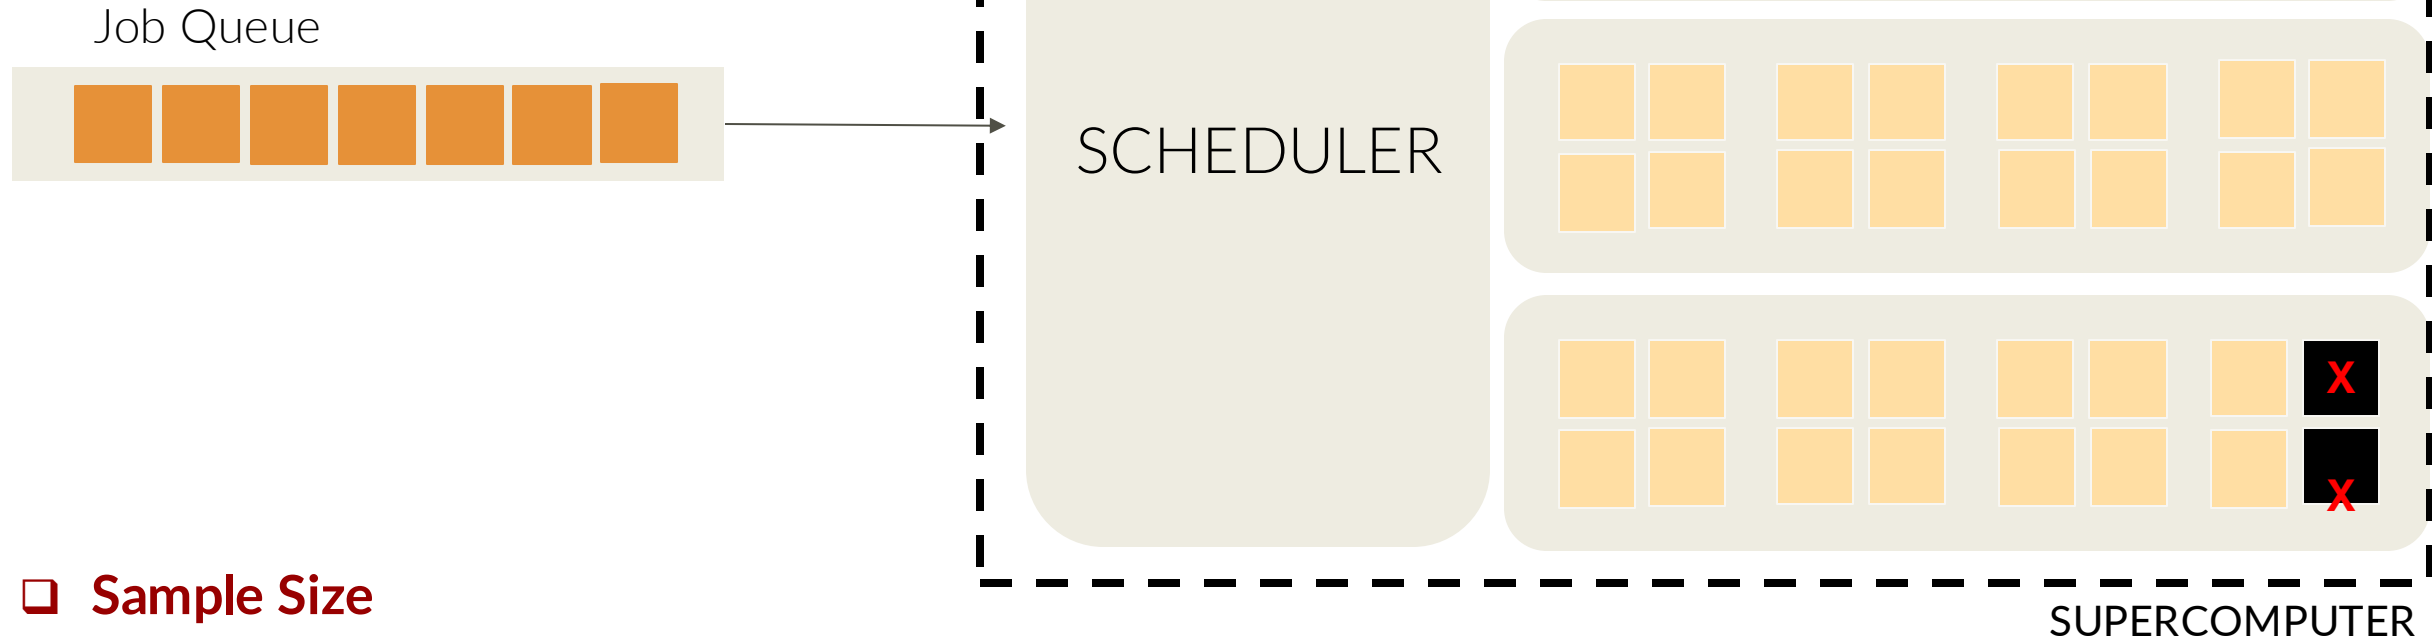
❑ Sample Size

- ❑ Sample measurements from almost all GPUs in each cluster
- ❑ Profiled 2.9x more GPUs than worst-case recommendations for statistical significance [Scogland, et al. SC'15]



Available Running SGEMM Unavailable Completed SGEMM

Methodology



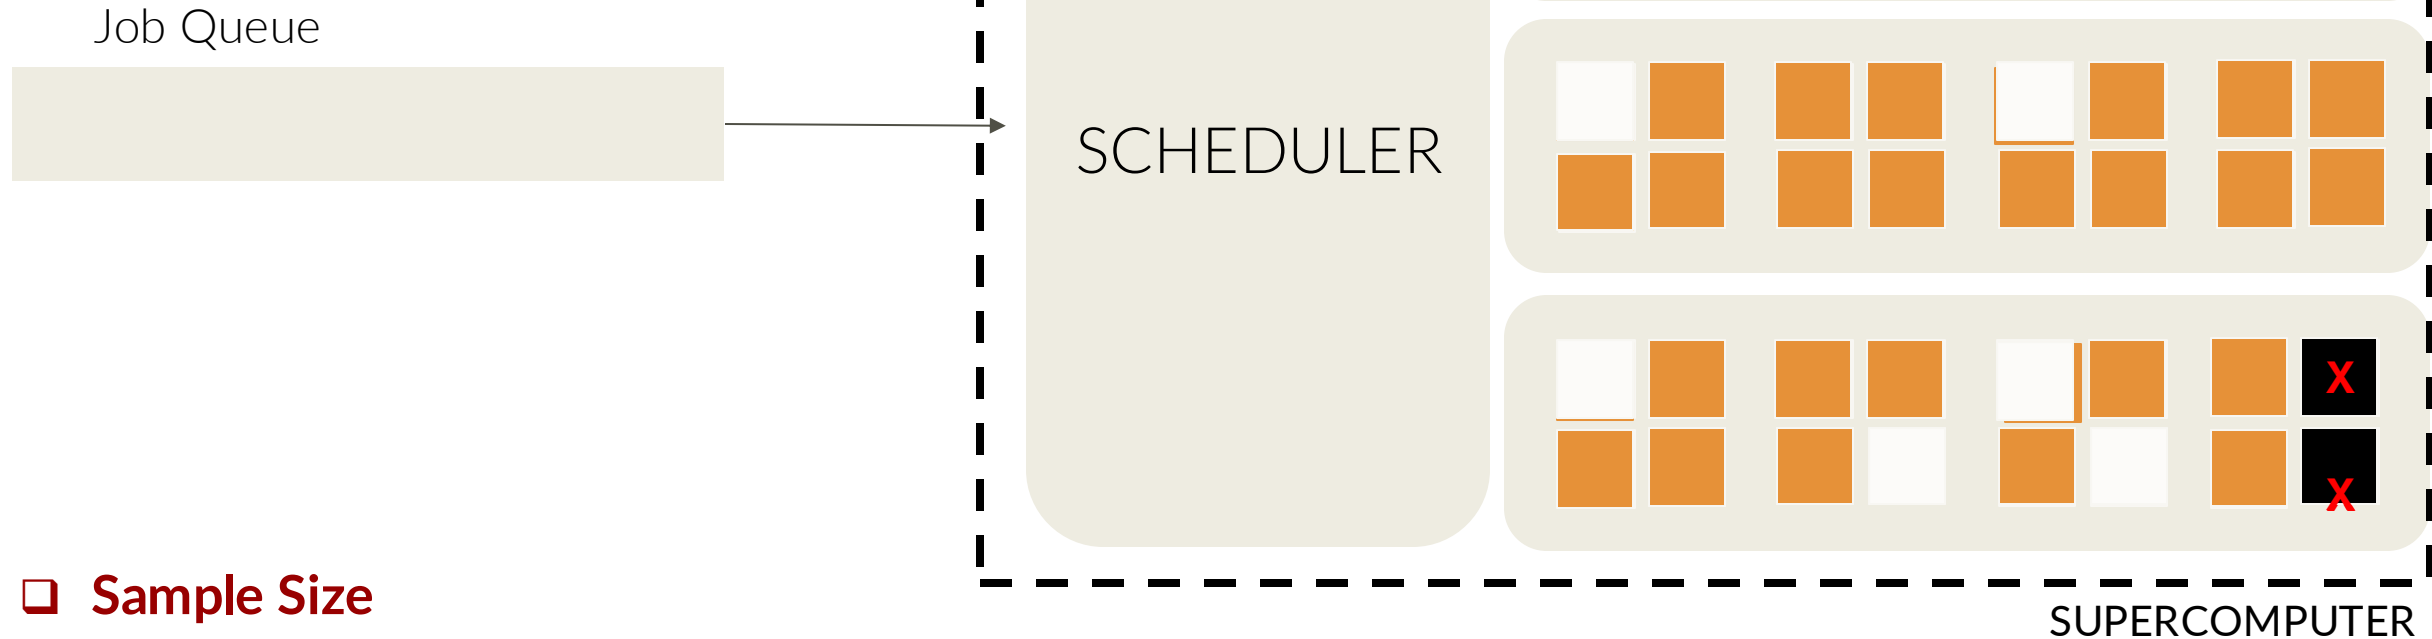
❑ Sample Size

- ❑ Sample measurements from almost all GPUs in each cluster
- ❑ Profiled 2.9x more GPUs than worst-case recommendations for statistical significance [Scogland, et al. SC'15]



Available Running SGEMM Unavailable Completed SGEMM

Methodology








❑ Sample Size

- ❑ Sample measurements from almost all GPUs in each cluster
- ❑ Profiled 2.9x more GPUs than worst-case recommendations for statistical significance [Scogland, et al. SC'15]



CLUSTER PARAMETERS



Cluster	GPU Architecture	Number of GPUs	Cooling Method
LLNL Corona	AMD MI60 	328	air cooled
TACC Longhorn	NVIDIA V100 	416	air cooled
TACC Frontera	NVIDIA Quadro RTX 5000 	360	mineral oil cooled
SNL Vortex	NVIDIA V100 	216	water cooled
ORNL Summit	NVIDIA V100 	27648	air cooled

Not All GPUs Are Created Equal: Characterizing Variability in Large-Scale, Accelerator-Rich Systems

<https://dl.acm.org/doi/abs/10.5555/3571885.3571971>



VARIABILITY ACROSS APPLICATIONS

